

Chapter 32. Class Odds and Ends

This chapter concludes our look at OOP in Python by presenting a collage of more advanced class topics. We will survey customizing built-in types, the relationship of classes and types, attribute tools like slots and properties, the special-case static and class methods, decorators and metaclasses, and the `super` call’s complete story. Some of these are introduced here but resumed by focused chapters in this book’s [Part VIII, “Advanced Topics”](#).

As we’ve seen, Python’s OOP model is, at its core, relatively simple, and some of the topics presented in this chapter are so advanced and optional that you may not encounter them very often in your Python applications-programming career. In the interest of completeness, though—and because you never know when an “advanced” topic may crop up in code you use—we’ll round out our discussion of classes with a brief look at these advanced tools for OOP work.

As usual, because this is the last chapter in this part of the book, it ends with a section on class-related “gotchas” and a set of lab exercises for this part to help cement the ideas we’ve studied here. Beyond these exercises, studying larger OOP Python projects or starting some of your own is heartily recommended as a supplement to this book. As with much in life and computing, the benefits of OOP tend to become more apparent with practice.

NOTE

Blast from the past: Python 3.X launched with a mandatory “*new-style*” class model that could be enabled in 2.X as an option; 2.X’s own model was dubbed “*classic*.¹ At least in terms of its OOP support, new-style classes transformed Python into a different language altogether—one that borrows much more from, and is often as complex as, other languages in this domain. The last chapter’s MRO and most topics in this chapter were part of this package. Because this book is now focused on 3.X only, the term “new style” is moot and unused here—all its classes qualify.

Extending Built-in Object Types

Besides implementing new kinds of objects, classes are sometimes used to extend the functionality of Python's built-in object types to support more exotic data structures. For instance, to add *queue* insert and delete methods to lists, you can code classes that wrap (embed) a list object and augment it with insert and delete methods that process the list specially, using the delegation technique we studied in [Chapter 31](#). You can also use simple inheritance to customize built-in types for such custom roles. The next two sections show both techniques in action.

Extending Types by Embedding

Do you remember those set functions we wrote in Chapters [16](#) and [18](#)? Here's what they look like brought back to life as a Python class. [Example 32-1](#) (file *setwrapper.py*) implements a new set object type by moving set functions to methods and adding some basic operator overloading. For the most part, this class just wraps a Python list with extra set operations. But because it's a class, it also supports multiple instances and customization by inheritance in subclasses. Unlike our earlier functions, using classes here allows us to make multiple self-contained set objects with preset data and behavior rather than passing lists into functions manually.

Example 32-1. *setwrapper.py*

```
class Set:
    def __init__(self, value = []):      # Constructor
        self.data = []                  # Manages a list
        self.concat(value)              # Removes duplicates

    def intersect(self, other):          # other is any iterable
        res = []                        # self is the subject
        for x in self.data:
            if x in other:              # Pick common items
                res.append(x)
        return Set(res)                # Return a new Set

    def union(self, other):             # other is any iterable
        res = self.data[:]             # Copy of my list
        for x in other:
            if not x in res:           # Add items in other
                res.append(x)
        return Set(res)
```

```

        res.append(x)
    return Set(res)

    def concat(self, value):          # value: list, Set
        for x in value:              # Removes duplicates
            if not x in self.data:
                self.data.append(x)

    def __len__(self):               return len(self.data)
    def __getitem__(self, key):      return self.data[key]
    def __and__(self, other):       return self.intersect(other)
    def __or__(self, other):        return self.union(other)
    def __repr__(self):             return f'Set({self.data!r})'
    def __iter__(self):             return iter(self.data)

```

To use this class, we make instances, call methods, and run defined operators as usual:

```

$ python3
>>> from setwrapper import Set
>>> x = Set([1, 3, 5, 7, 3])
>>> x.union(Set([1, 4, 7]))
Set([1, 3, 5, 7, 4])
>>> x | Set([1, 4, 6, 4])
Set([1, 3, 5, 7, 4, 6])

```

Overloading operations such as indexing and iteration also enables instances of our `Set` class to often masquerade as real lists. Because you will interact with and extend this class in an exercise at the end of this chapter, we'll put this code on the back burner until its solution in [Appendix B](#).

Extending Types by Subclassing

While the prior section's embedding works, Python's built-in types can also be subclassed directly. In fact, type-conversion functions such as `list`, `str`, `dict`, and `tuple` are really built-in type names; although transparent to your script, a type-conversion call (e.g., `list('text')`) is really an invocation of a type's constructor.

This allows you to customize or extend the behavior of built-in types with user-defined `class` statements: simply subclass the type names to customize

them. Instances of your type subclasses can generally be used anywhere that the original built-in type can appear. For example, suppose you have trouble getting used to the fact that Python list offsets begin at 0 instead of 1. Not to worry—you can always code your own subclass that customizes this core behavior of lists, and [Example 32-2](#) shows how.

Example 32-2. typesubclass.py

```
"""
Subclass built-in list type/class.
Map 1..N to 0..N-1, call back to built-in version.
"""

class MyList(list):
    def __getitem__(self, offset):
        print(f'<indexing {self} at {offset}>')
        return list.__getitem__(self, offset - 1)

if __name__ == '__main__':
    print(list('abc'))
    x = MyList('abc')                      # __init__ inheritance
    print(x)                                # __str__ / __repr__

    print(x[1])                            # MyList.__getitem__
    print(x[3])                            # Customizes list's __getitem__

    x.append('hack!'); print(x)            # Attributes from list
    x.reverse(); print(x)
```

In this file, the `MyList` subclass extends the built-in list's `__getitem__` indexing method only, to map indexes 1 to N back to the required 0 to N-1. Really, all it does is decrement the submitted index and call back to the superclass's version of indexing, but it's enough to do the trick:

```
$ python3 typesubclass.py
['a', 'b', 'c']
['a', 'b', 'c']
<indexing ['a', 'b', 'c'] at 1>
a
<indexing ['a', 'b', 'c'] at 3>
c
```

```
['a', 'b', 'c', 'hack!']
['hack!', 'c', 'b', 'a']
```

This output also includes tracing text the class prints on indexing. Of course, whether changing indexing this way is a good idea, in general, is *another issue*—users of your `MyList` class may very well be confused by such a core departure from Python sequence behavior. The ability to customize built-in types this way can be a powerful asset, though.

For instance, this coding pattern gives rise to an alternative way to code a set—as a subclass of the built-in list type rather than a standalone class that manages an embedded list object, as shown in the prior section. As discussed in [Chapter 5](#), Python today comes with a powerful built-in set object, along with literal and comprehension syntax for making new sets. Coding one yourself, though, is still a great way to learn about type subclassing in general.

The code in [Example 32-3](#), file `setsubclass.py`, customizes lists to add just methods and operators related to set processing. Because all other behavior is inherited from the built-in `list` superclass, this makes for a shorter and simpler alternative—everything not defined here is routed to `list` directly.

Example 32-3. `setsubclass.py`

```
class Set(list):
    def __init__(self, value = []):          # Constructor
        list.__init__(self)                  # Customizes l
        self.concat(value)                 # Copies mutat

    def intersect(self, other):             # other is any
        res = []                           # self is the
        for x in self:
            if x in other:                # Pick common
                res.append(x)
        return Set(res)                  # Return a new

    def union(self, other):               # other is any
        res = Set(self)                  # Copy me and
        res.concat(other)
        return res

    def concat(self, value):              # value: list,
        for x in value:                 # Removes dupl
            if not x in self:
```

```

        self.append(x)

    def __and__(self, other): return self.intersect(other)
    def __or__(self, other): return self.union(other)
    def __repr__(self): return f'Set({list.__repr__(self)})'

if __name__ == '__main__':
    x = Set([1, 3, 5, 7])
    y = Set([2, 1, 4, 5, 6])
    print(x, y, len(x))
    print(x.intersect(y), y.union(x))
    print(x & y, x | y)
    x.reverse(); print(x)

```

Here is the output of the self-test code at the end of this file. Because subclassing core types is a somewhat advanced feature with a limited audience, we'll end this topic here, but you're invited to trace through these results in the code to study its behavior:

```

$ python3 setsubclass.py
Set([1, 3, 5, 7]) Set([2, 1, 4, 5, 6]) 4
Set([1, 5]) Set([2, 1, 4, 5, 6, 3, 7])
Set([1, 5]) Set([1, 3, 5, 7, 2, 4, 6])
Set([7, 5, 3, 1])

```

Subtleties: some inherited list operations may introduce duplicates to our `Set`, and there are more efficient ways to implement sets with dictionaries in Python, which replace the nested linear search scans in the set implementations shown here with more direct dictionary index operations (hashing) and so run much quicker. If you're interested in sets, also take another look at the `set` object type we explored in [Chapter 5](#); this type provides extensive set operations as built-in tools. Set implementations are fun to experiment with but not strictly required in Python today.

More important here is the question of why we can subclass built-in types like `list` at all. The next section solves the mystery—at least as much as this chapter can.

The Python Object Model

The reason we could subclass built-in types in the prior section is that types and classes are largely one and the same—a unification that came with the “new-style” model alluded to at the start of this chapter. For built-ins, some instances can uniquely be coded with literal syntax like `[]`, `'1p6e'`, and `3.12` instead of class calls like `list()`, `str()`, and `float()`, but they are instances of a class, nonetheless.

In fact, built-in types and user-defined classes are both classes and are both themselves instances of the built-in `type` class. The `type` object generates classes as its instances, classes generate instances of themselves, and classes are really just user-defined types. And on top of all this, the built-in `object` class provides defaults for every object.

Classes Are Types Are Classes

While you probably shouldn’t ponder the preceding definitions before operating heavy machinery, it’s easy to see all this in code. The `type` built-in with one argument returns any object’s type, which is normally the same as the object’s `__class__`, and `isinstance` checks whether an object inherits from another. Here’s the story for user-defined *classes* (a.k.a. types):

```
>>> class Hack: pass          # A humble user-defined class
>>> I = Hack()                # Make an instance by calling the class

>>> type(I)                  # Type is the user-defined class
<class '__main__.Hack'>

>>> type(Hack)                # User-defined classes are classes
<class 'type'>

>>> I.__class__, Hack.__class__
(<class '__main__.Hack'>, <class 'type'>)

>>> isinstance(I, object), isinstance(Hack, object)
(True, True)
```

This works the same for built-in *types* (a.k.a. classes), but there is also literal syntax for generating instances:

```
>>> I = 'hack'                      # Make an instance by literal assignment
>>> type(I)                         # Built-in objects are instances
<class 'str'>

>>> type(str)                       # Built-in classes are instances
<class 'type'>

>>> I.__class__, str.__class__
(<class 'str'>, <class 'type'>)

>>> isinstance(I, object), isinstance(str, object)
(True, True)
```

In fact, `type` itself reports in as a class, though it has no type but itself—circularly capping the chain:

```
>>> type(type)                      # The type class ends the chain
<class 'type'>
>>> type(type(type))                # Hmm...the top of the chain
<class 'type'>
```

This model may seem academic (and to some extent is), but it allows us to specialize built-in types with normal user-defined classes and bears on type-testing code: you must know what a type is to test it accurately.

Some Instances Are More Equal Than Others

It's tempting to simply take away from the foregoing that classes and types are the same, but this story is richer than that may imply: the instances we make from these objects diverge in both functionality and inheritance.

To truly understand how, we have to briefly factor in *metaclasses*—classes that generate other classes. The `type` built-in itself *is* a metaclass and may be *customized* with user-defined subclasses. These subclasses are coded with normal `class` statements, selected with special syntax in `class` headers, and designed to play metaclass roles, but they won't be covered in full until [Chapter 40](#).

In brief, though, the complete relationship between instances, classes, and types is as follows:

- *Instances* are created from classes—both built-in and user-defined.
- *Classes* themselves are created from the built-in `type` class or one of its subclasses.
- The `object` built-in class is a superclass to every object—instance, class, or both.

Although everything is ultimately an “instance” in Python, there are two fundamentally different *kinds* of instances, and conflating these only serves to mask the true complexity of the model. It doesn’t help to distinguish these as instances of *user-defined* classes or not: metaclasses may be user-defined classes too. Nor is this about being a *subclass* of a type: the real fork in this model is that classes are *created* from a type class specially.

As you’ll learn in full later, classes define their types with optional `metaclass` syntax that defaults to `type` if omitted, but other instances define their types by the class calls or literal syntax we’ve used so far:

```
class C(metaclass=Meta): ...      # Class creation: metaclass  
I = C(...)                      # Instance creation: use metaclass  
X = [1, 2]                        # Instance creation: built-in
```

While both produce instances in some sense, these different syntaxes create very different kinds of instances—*class* and *nonclass*—with fundamentally different behaviors:

Nonclass instances do not make instances

Classes are created from `type` (or another metaclass), similar to the way instances are created from classes. Once created, though, the analogy fails: *classes* create instances of their own, but *nonclass* instances do not.

Classes have an extra inheritance search

There are really *two* inheritance trees and searches in Python, which are distinct but not entirely disjoint. The secondary tree is formed by `type` and its subclasses and is searched only for *classes*, not *nonclass* instances.

In other words, nonclass instances seal off the instantiation chain, and inheritance differs for nonclass instances and classes themselves—even though the latter are also instances of `type`. All of this boils down to different creation syntax that makes different kinds of objects, which are often confusingly lumped together as “instances”:

```
>>> class C: pass                      # A ty
>>> I = C()                           # A nc
>>> isinstance(I, type), isinstance(C, type)  # Only
(False, True)
```

While types and classes may be synonymous, the instances we create from them vary per creation code.

The Inheritance Bifurcation

Though we can’t get into full details here, the inheritance search used for *classes* (a.k.a. types) differs from what we’ve seen so far and may be their most profound distinction. In short, inheritance is always based on the *MRO* (method resolution order) we studied in [“Multiple Inheritance and the MRO”](#), but varies as follows:

Nonclass inheritance

As we’ve seen, inheritance run on a *nonclass* instance searches the `__dict__` attributes of instance, class, and superclasses, per the MRO order we studied in the prior chapter. This works by first checking the instance, then following the instance’s `__class__` to its class, and finally following each class’s `__bases__` to superclasses. Technically, `__bases__` are used to make an `__mro__` at `__class__`, which inheritance scans.

Class inheritance

Inheritance run on a *class* directly, though, first searches the `__dict__` of the class and all its supers available from `__bases__` as usual, but then *also* searches the separate class tree formed by the `type` class and its *metaclass* subclasses. The second part of this works by following the class’s own `__class__` to its `type` class tree and using `__bases__` and MROs there, too—but only as a last resort and only for inheritance run on classes.

In fact, if you know where to look, you can inspect the inheritance sources that differ for nonclass instances like `I` and classes like `C` in the prior example—though the underscores and displays aren’t pretty:

```
>>> isinstance(I, C), type(I)
(True, <class '__main__.C'>)

>>> I, I.__class__, I.__class__.__bases__
(<__main__.C object at 0x101265d60>, <class '__main__.C'>)

>>> C, C.__bases__, C.__class__, C.__class__.__bases__
(<class '__main__.C'>, (<class 'object'>,), <class 'typ...
```

But due to the way MROs are computed from `__bases__` and scanned, it’s more accurate to think of inheritance’s different search orders for nonclass instances and classes as follows—where each `__mro__` is a *flattened tree*:

```
>>> I, I.__class__.__mro__
(<__main__.C object at 0x101265d60>, (<class '__main__.C'>))

>>> C.__mro__, C.__class__.__mro__
((<class '__main__.C'>, <class 'object'>), <class 'typ...
```

And because each item’s `__dict__` is checked, the ordered set of candidates searched by inheritance for nonclass instances `I` and classes `C` is ultimately and respectively as follows—with *two* MRO scans of flattened trees for classes only, and ignoring the fact that some kinds of *descriptors*, introduced ahead, take precedence in both trees as you’ll learn in [Chapter 40](#):

```
[I.__dict__] + [x.__dict__ for x in I.__class__.__mro__
[x.__dict__ for x in C.__mro__] + [x.__dict__ for x in
```

Wait—there’s a *second* tree in inheritance? Well, yes, though it doesn’t come into play in the vast majority of application code. The type/metaclass tree is used in advanced class-management roles and, even then, is often limited to class customization at class creation time.

Still, this secondary tree, along with the descriptors' special cases omitted here, bifurcates and convolutes the inheritance story, especially compared to its prior forms. It also explains why some class attributes like `__bases__` are not inherited by nonclass instances—they're located in the secondary tree (i.e., MRO) searched only for classes:

```
>>> '__bases__' in I.__dict__          # Not in instance dict
False
>>> '__bases__' in C.__dict__          # Not in class dict
False
>>> '__bases__' in C.__class__.__dict__ # In metaclass dict
True
```

Because of the two-tree inheritance model, such names inherited by classes are not inherited by their instances:

```
>>> C.__bases__                      # Instance has no __bases__
(<class 'object'>,)
>>> I.__bases__
AttributeError: 'C' object has no attribute '__bases__'
```

The Metaclass/Class Dichotomy

So, where does this odd tale of type/class unification leave us? Types indeed behave as classes, and this allows us to extend them with normal class syntax in both the primary and metaclass trees. But it also comes with noticeable *seams*, including special-case syntax for class instantiation, an extra type-tree search for classes only, two very different kinds of instances, and unique semantics for metaclasses that customize types (a.k.a. classes), which we'll uncover later.

In fact, a reasonable argument can be made that the *type/class* dichotomy of earlier Pythons may simply have morphed into one of *metaclass/class*—which trades a straightforward distinction for all the seams just enumerated and muddles inheritance and the fundamental meaning of *names* in Python everywhere to support what in the end is a very rare use case. As usual, the net merit of the morph is yours to weigh.

To be fair, some of the widespread confusion this model has spawned may stem from `type` itself: it's overloaded to either *return* a sole argument's type, or *generate* a new instance of itself for multiple arguments—just like other constructors and equivalent to what a `class` statement does to make a class object:

```
type(object)                      # Fetches the type of an object
type(classname, superclasses, attributedict)  # Makes a new class object
```

The first of these roles might have been better named “`typeof`,” but the second will have to await the metaclass preview later in this chapter and the extended coverage in [Chapter 40](#).

And One “object” to Rule Them All

To round out this topic, keep in mind that because topmost classes inherit from the built-in class `object`, every object *derives* (i.e., inherits) from it, whether directly or through a superclass—and whether you code `object` or not:

```
>>> class C: pass
>>> class D(object): pass

>>> dir(C) == dir(D)
True
>>> C.__bases__, D.__bases__
((<class 'object'>,), (<class 'object'>,))
```

In fact, the `type` class inherits from the `object` class, and `object` inherits from `type`, even though the two are different objects—a circular relationship that crowns the object model and may make your cranium catch fire (to avoid combustion, keep in mind that `isinstance` is true for *either* a subclass relationship or creation source, though this is based on inheritance through the secondary type-class tree for classes):

```
>>> type is object
False
>>> type(type), type(object)
(<class 'type'>, <class 'type'>)
```

```
>>> isinstance(type, object), isinstance(object, type)
(True, True)
>>> type.__bases__, object.__bases__
((<class 'object'>,), ())
```

Strange though it may seem, this has a number of practical consequences. For one thing, it means that we sometimes must be aware of the method defaults that come with the implicit (or explicit) `object` root class. As we noted in earlier chapters, for instance, the `object` class comes with a `__repr__` for display:

```
>>> class C: pass                      # All classes inherit
>>> X = C()
>>> X.__repr__
<method-wrapper '__repr__' of C object at 0x1091a7920>
```

For another, this also allows us to write code that can safely assume and use an `object` superclass. As an example, we can rely on it to be a call-chain “anchor” in some `super` built-in roles described ahead and can reroute method calls to it from attribute-interceptor methods to invoke higher default behavior. Per [Chapter 30](#):

```
object.__setattr__(self, attr, value)
```

We’ll code examples of such rerouting later in the book; for now, let’s move on to something a bit more tangible and our next topic in this OOP jamboree.

Advanced Attribute Tools

Along with the normal class and instance attributes we’ve been using so far, Python’s OOP support includes attribute tools of narrower scope—*slots*, *properties*, *descriptors*, and more. Slots, for example, are an optimization option, and properties and descriptors allow classes to augment access. None of these tools are required, but as for most topics in this chapter, all are fair game in Python code you may someday use. Most of these tools get extended coverage in [Chapter 38](#), but slots get full coverage here, and others are presented in abbreviated form.

Slots: Attribute Declarations

First off, we've noted the implications of *slots* several times in this part of the book. In short, by assigning an iterable of attribute name strings to a special `__slots__` class attribute, we can enable a class to both limit the set of legal attributes that instances of the class will have and optimize memory usage and possibly program speed. As you'll find, though, slots should be used only in applications that clearly warrant the added complexity. They will complicate your code, may complicate or break code you may use, and rigidly require universal deployment to be effective.

Slot basics

To declare slots, assign an iterable (e.g., list) of string names to the special `__slots__` variable and attribute at the top level of a `class` statement: only those names in `__slots__` can be assigned as instance attributes. This doesn't change the way these attributes work in general, though; like all names in Python, instance attribute names must always be assigned before they can be referenced, even if they're listed in `__slots__`. Here are the basics:

```
>>> class Limiter(object):
...     __slots__ = ['age', 'name', 'job']

>>> I = Limiter()
>>> I.age
AttributeError: 'Limiter' object has no attribute 'age'

>>> I.age = 40
>>> I.age
40
>>> I.ape = 1000
AttributeError: 'Limiter' object has no attribute 'ape'
```

This feature is advertised as both a way to catch typo errors like this (assignments to illegal attribute names not in `__slots__` are detected instead of silently assigned), as well as an optimization mechanism that saves memory.

Allocating a namespace dictionary for every instance object can be expensive in terms of *memory* if many instances are created and only a few attributes are required. To save space, instead of allocating a dictionary for each instance, Python reserves just enough space in each *instance* to hold a value for each slot attribute, along with inherited attributes in the common *class* to manage slot access. This might additionally *speed* execution, though this benefit may vary per program, platform, and Python version (spoiler: the speedup is trivial today, as we'll prove ahead).

You shouldn't normally use slots

Slots are a fairly major break with Python's core dynamic nature, which dictates that any name may be created by assignment (and frankly, tend to appeal most to people with backgrounds in draconian languages). In fact, they partly imitate C++ for efficiency at the expense of flexibility and even have the potential to *break* some programs.

As you'll see, slots also come with a plethora of special-case usage *rules*. Per Python's own manual, they should *not* be used except in clearly warranted cases—they are difficult to deploy correctly, complicate your code badly, and are best limited to very rare memory-critical programs that produce an extremely large numbers of instances.

In other words, this is yet another feature that should be used only if clearly justified. Unfortunately, slots seem to be showing up in Python code much more often than they should; their obscurity seems to be a draw in itself. Slots are actually used by Python, unlike *type hinting*, their declaration cousin of [Chapter 6](#), but they are similarly paradoxical and restrictive. As usual, knowledge is your best ally in such things, so let's take a deeper look here.

Slots and namespace dictionaries

Potential benefits aside, slots can complicate a class model—and code that relies on it—substantially. In fact, some instances with slots may not have a `__dict__` attribute namespace dictionary at all, and others will have data attributes that this dictionary does not include. To be clear, this is a *major incompatibility* with the traditional class model—one that can impact any code that accesses attributes generically and may even cause some to fail altogether.

For instance, programs that list or access instance attributes by name string may need to use more storage-neutral interfaces than `__dict__` if slots may be used. Because an instance's data may include class-level names such as slots—either in *addition* to or *instead* of namespace dictionary storage—both attribute sources may need to be queried for completeness, and some roles may be rendered impossible.

Let's see what this means in terms of code and explore more about slots along the way. First off, when slots are used, instances do not normally have an attribute dictionary—instead, Python uses the class *descriptors* feature introduced ahead to allocate and manage space reserved for slot attributes in the instance:

```
>>> class C:  
    __slots__ = ['a', 'b']          # __slots__ me  
  
>>> I = C()  
>>> I.a = 1  
>>> I.a  
1  
>>> I.__dict__  
AttributeError: 'C' object has no attribute '__dict__'.  

```

However, we can still fetch and set slot-based attributes by name string using storage-neutral tools such as `getattr` and `setattr` (which look beyond the instance `__dict__` and thus include class-level names like slots) and list them with `dir` (which collects all inherited names of any kind throughout a class tree):

```
>>> getattr(I, 'a')  
1  
>>> setattr(I, 'b', 2)          # But getattr()  
>>> I.b  
2  
>>> 'a' in dir(I)            # And dir() fi  
True  
>>> 'b' in dir(I)            # Though __dic  
True  
>>> I.__dict__  
AttributeError: 'C' object has no attribute '__dict__'.  

```

Also keep in mind that without an attribute namespace dictionary, it's not possible to assign *new* names to instances that are not names in the slots list:

```
>>> class D:  
    __slots__ = ['a', 'b']  
    def __init__(self):  
        self.d = 4                      # Cannot add r  
  
>>> I = D()  
AttributeError: 'D' object has no attribute 'd'
```

We can still accommodate extra attributes, though, by including `__dict__` explicitly in `__slots__` in order to create an attribute namespace dictionary in *addition* to slots:

```
>>> class D:  
    __slots__ = ['a', 'b', '__dict__']      # Name __  
    c = 3                                # Class c  
    def __init__(self):  
        self.d = 4                      # d stored  
  
>>> I = D()  
>>> I.d  
4  
>>> I.c  
3  
>>> I.a                      # All instance attrs in __  
AttributeError: 'D' object has no attribute 'a'  
>>> I.a = 1  
>>> I.b = 2
```

In this case, *both* storage mechanisms are used. This renders `__dict__` too limited for code that wishes to treat slots as instance data, but generic tools such as `getattr` still allow us to process both storage forms as a single set of attributes:

```
>>> I.__dict__                  # Some objects have bc  
{'d': 4}                         # getattr() can fetch  
>>> I.__slots__  
['a', 'b', '__dict__']
```

```
>>> getattr(I, 'a'), getattr(I, 'c'), getattr(I, 'd')
(1, 3, 4)
```

Because `dir` also returns all *inherited* attributes, though, it might be too broad in some contexts; it also includes class-level methods and even all `object` defaults. Code that wishes to list *just* instance attributes may, in principle, still need to allow for both storage forms explicitly. We might at first naively code this as follows:

```
>>> for attr in list(I.__dict__) + I.__slots__:
    print(attr, '=>', getattr(I, attr))
```

Since either can be omitted, we may more correctly code this as follows, using `getattr` to allow for defaults—a noble but nonetheless inaccurate approach, as the next section will explain:

```
>>> for attr in list(getattr(I, '__dict__', [])) + geta
    print(attr, '=>', getattr(I, attr))

d => 4
a => 1
b => 2
__dict__ => {'d': 4}
```

Multiple `__slot__` lists in superclasses

The preceding code works in this specific case, but in general, it's not entirely accurate. Specifically, this code addresses only slot names in the *lowest* `__slots__` attribute inherited by an instance, but slot lists may appear more than once in a class tree. That is, a name's absence in the lowest `__slots__` list does not preclude its existence in a higher `__slots__`. Because slot names become class-level attributes, instances acquire the *union* of all slot names anywhere in the tree by the normal inheritance rule:

```
>>> class E:
    __slots__ = ['c', 'd']                      # Superclass

>>> class D(E):
    __slots__ = ['a', '__dict__']      # But so does
```

```

>>> I = D()                                     # The instance
>>> dir(I)
[...names omitted..., 'a', 'c', 'd']
>>> I.a = 1; I.b = 2; I.c = 3                  # slots: a, c
>>> I.a, I.c
(1, 3)

```

But inspecting just the inherited slots list won't pick up slots defined *higher* in a class tree:

```

>>> E.__slots__                                # But __slots__
['c', 'd']
>>> D.__slots__
['a', '__dict__']
>>> I.__slots__                               # Instance inherits
['a', '__dict__']
>>> I.__dict__                                 # And has its own
{'b': 2}

>>> for attr in list(getattr(I, '__dict__', [])) + geta
      print(attr, '=>', getattr(I, attr))

b => 2                                         # Other superclasses
a => 1
__dict__ => {'b': 2}

>>> dir(I)                                    # But dir() includes
[...names omitted..., 'a', 'b', 'c', 'd']

```

In other words, in terms of listing instance attributes generically, one `__slots__` isn't always enough—they are potentially subject to the full inheritance search procedure. If multiple classes in a class tree may have their own `__slots__` attributes, tools must develop other policies for listing attributes—as the next section explains.

Handling slots and other “virtual” attributes generically

The prior chapter concluded with a brief summary of the slots policies of its attribute lister tools—a prime example of why generic programs may need to care about slots. Such tools that attempt to list instance data attributes generically must account for slots and perhaps other such “virtual” instance

attributes like *properties* and *descriptors* introduced ahead—names that similarly reside in classes but may provide attribute values for instances on request. Slots are the most data-centric of these but are representative of a larger category.

Such attributes require inclusive approaches, special handling, or general avoidance—the latter of which becomes unsatisfactory as soon as any programmer uses slots in subject code. Really, class-level instance attributes like slots probably necessitate a redefinition of the term *instance data*—as locally stored attributes, the union of all inherited attributes, or some subset thereof.

For example, some programs might classify slot names as attributes of *classes* instead of instances; these attributes do not exist in instance namespace dictionaries, after all. Alternatively, as shown earlier, programs can be more inclusive by relying on `dir` to fetch all inherited attribute names and `getattr` to fetch their corresponding values—without regard to their physical location or implementation. If you must support slots as instance data, this may be the most robust way to proceed:

```
>>> class Slotful:
...     __slots__ = ['a', 'b', '__dict__']
...     def __init__(self, data):
...         self.c = data
...
>>> I = Slotful(3)
>>> I.a, I.b = 1, 2
>>> I.a, I.b, I.c                                # Normal __c
(1, 2, 3)
...
>>> I.__dict__                                     # Both __c
{'c': 3}
>>> [x for x in dir(I) if not x.startswith('__')]
['a', 'b', 'c']
...
>>> I.__dict__['c']                               # __dict__
3
>>> getattr(I, 'c'), getattr(I, 'a')            # dir+getattr
(3, 1)                                         # applies
...
>>> for a in (x for x in dir(I) if not x.startswith('__'))
...     print(a, '=>', getattr(I, a))
```

- a 1
- b 2
- c 3

Under this `dir` / `getattr` model, you can still map attributes to their inheritance sources and filter them more selectively by source or type, if needed, by scanning the *MRO*—as we did in the prior chapter’s *mapattrs.py* ([Example 31-14](#)). As a bonus, such tools and policies for handling slots will potentially apply automatically to properties and descriptors too, though these attributes are more explicitly computed values, and less obviously instance-related data than slots.

Also keep in mind that this is not just a tools issue. Class-based instance attributes like slots also impact the traditional coding of the `__setattr__` operator-overloading method we met in [Chapter 30](#). Because slots and some other attributes are not stored in the instance `__dict__`, and may even imply its *absence*, classes must instead generally run attribute assignments by rerouting them to the `object` superclass.

Slot usage rules

Slot declarations can appear in multiple classes in a class tree, but when they do, they are subject to a number of constraints that are somewhat difficult to rationalize unless you understand the implementation of slots as class-level *descriptors* for each slot name that are inherited by the instances in which the managed space is reserved (again, you’ll meet descriptors briefly ahead). Here are the main constraints that slots impose:

- **Slots in subs are pointless when absent in supers.** If a subclass inherits from a superclass without a `__slots__`, the instance `__dict__` attribute created for the superclass will always be accessible, making a `__slots__` in the subclass largely pointless. The subclass still manages its slots but doesn’t compute their values in any way and doesn’t avoid a dictionary—the main reason to use slots.
- **Slots in supers are pointless when absent in subs.** Similarly, because the meaning of a `__slots__` declaration is limited to the class in which it appears, subclasses will produce an instance `__dict__` if they do not define a `__slots__`, rendering a `__slots__` in a superclass largely pointless.

- **Redefinition renders super slots pointless.** If a class defines the same slot name as a superclass, its redefinition hides the slot in the superclass per normal inheritance. You can access the version of the name defined by the superclass slot only by fetching its descriptor directly from the superclass.
- **Slots prevent class-level defaults.** Because slots are implemented as class-level descriptors (along with per-instance space), you cannot use class attributes of the same name to provide defaults as you can for normal instance attributes: assigning the same name in the class overwrites the slot descriptor.
- **Slots cannot be combined in multiple inheritance.** Multiple inheritance cannot be used if more than one of the classes mixed together have nonempty slots lists—even if their slots define the same names. You’ll get an error when running the class that does the mixing. Empty slots lists allow the mixer to define slots or not, as desired.
- **Slots can impact `__dict__`.** As shown earlier, `__slots__` preclude both an instance `__dict__` and assigning names not listed, unless `__dict__` is listed explicitly too. Slots similarly preclude a `__weakref__` attribute used to support instance “weak references” covered briefly in [Chapter 6](#), but these are rare enough to soft-pedal here.

We’ve already seen the last of these in action. It’s easy to demonstrate how the new rules here translate to actual code—most crucially, a namespace dictionary is created when any class in a tree omits slots, thereby negating the memory optimization benefit but also supporting classes that require a `__dict__` when mixed in with others:

```

>>> class C: pass                                # Bullet 1: s]
>>> class D(C): __slots__ = ['a']                # Makes instar
>>> I = D()                                     # But slot nan
>>> I.a = 1; I.b = 2
>>> I.__dict__
{'b': 2}
>>> D.__dict__.keys()
dict_keys([... '__slots__', 'a', ...])

>>> class C: __slots__ = ['a']                   # Bullet 2: s]
>>> class D(C): pass                          # Makes instar
>>> I = D()                                     # But slot nan
>>> I.a = 1; I.b = 2
>>> I.__dict__

```

```

{'b': 2}

>>> C.__dict__.keys()
dict_keys([... '__slots__', 'a', ...])

>>> class C: __slots__ = ['a']                      # Bullet 3: or
>>> class D(C): __slots__ = ['a']                  # Superclass s

>>> class C: __slots__ = ['a']; a = 99    # Bullet 4: nc
ValueError: 'a' in __slots__ conflicts with class variab

>>> class C: __slots__ = ['a']                      # Bullet 5: or
>>> class D: __slots__ = ['a']                  # Use empty sl
>>> class E(C, D): pass
TypeError: multiple bases have instance lay-out conflic

```

In other words, besides their program-breaking potential, slots essentially require both *universal* and *careful* deployment to be effective—because slots do not compute values dynamically like properties (coming up in the next section), they are largely pointless unless each class in a tree uses them and is cautious to define only new slot names not defined by other classes. It’s an *all-or-nothing* feature—an unfortunate property shared by the `super` call discussed ahead:

```

>>> class C: __slots__ = ['a']                      # Assumes universal
>>> class D(C): __slots__ = ['b']                  # And may break
>>> I = D()
>>> I.a = 1; I.b = 2
>>> I.__dict__
AttributeError: 'D' object has no attribute '__dict__'.
>>> C.__dict__.keys(), D.__dict__.keys()
(dict_keys([... '__slots__', 'a', ...]), dict_keys([... '__s'])

```

Such rules—and others omitted here for space—are part of the reason slots are not widely used and are not generally recommended except in pathological cases where their space reduction is significant. Even then, their potential to complicate or break code should be ample cause to carefully consider the trade-offs. Not only must they be spread almost *neurotically* throughout a framework, but they may also break tools you rely on.

Example impacts of slots: ListTree and mapattrs

As a more realistic example of slots' effects, due to the first bullet in the prior section, [Chapter 31](#)'s `ListTree` class ([Example 31-13](#)) does *not fail* when mixed into a class that defines `__slots__`, even though it scans instance namespace dictionaries without verifying their presence. This lister class's own lack of slots is enough to ensure that the instance will still have a `__dict__` and hence not trigger an exception when fetched or indexed.

For example, both of the following *single*-inheritance trees display without error—the second also allows names not in the slots list to be assigned as instances attributes, including any required by the superclass:

```
class C(ListTree): pass
I = C()                                     # OK: nc
print(I)

class C(ListTree): __slots__ = ['a', 'b']      # OK: su
I = C()
I.c = 3
print(I)                                     # Displa
```

The following *multiple*-inheritance classes display correctly as well—*any* nonslot class like `ListTree` generates an instance `__dict__` and can thus safely assume its presence. Although it renders subclass slots pointless, this is a positive side effect for tool classes like `ListTree` and its [Chapter 28](#) predecessor:

```
class A: __slots__ = ['a']                      # Both C
class B(A, ListTree): pass
print(B())

class A: __slots__ = ['a']
class B(A, ListTree): __slots__ = ['b']          # Displa
print(B())
```

In general, though, tools may need to catch exceptions when `__dict__` is absent or use a `hasattr` or `getattr` to test or provide defaults if slot usage may preclude an instance namespace dictionary. For instance, [Chapter 31](#)'s `mapattrs.py` module ([Example 31-14](#)) must check for

`__dict__` presence explicitly because it is not a class mixed into others, and so cannot assume this attribute. Like `ListTree`, this example also associates slots with their classes.

Run these examples on your own for more info. Slots' impacts may be onerous, but knowledge is your best defense.

What about slots speed?

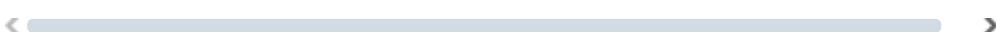
Finally, while slots primarily optimize memory use, their speed impact is less clear-cut. [Example 32-4](#) codes a simple test script using the `timeit` techniques we studied in [Chapter 21](#). For both the slots and nonslots (instance dictionary) storage models, it makes 1,000 instances, assigns and fetches 4 attributes on each, and repeats 1,000 times—for both models taking the best of 5 runs that each exercise a total of 8M attribute operations.

Example 32-4. slots-test.py

```
import timeit
base = """
Is = []
for i in range(1000):
    I = C()
    I.a = 1; I.b = 2; I.c = 3; I.d = 4
    t = I.a + I.b + I.c + I.d
    Is.append(I)
"""

stmt = """
class C:
    __slots__ = ['a', 'b', 'c', 'd']
""" + base
print('Slots    =>', end=' ')
print(min(timeit.repeat(stmt, number=1000, repeat=5)))

stmt = """
class C:
    pass
""" + base
print('Nonslots=>', end=' ')
print(min(timeit.repeat(stmt, number=1000, repeat=5)))
```



At least for this code, on the macOS test host, and using CPython 3.12, the best times imply that slots are only slightly quicker, though this says little about memory space and is prone to change arbitrarily in the future (PyPy 7.3 struggled on this test with times 10x slower than CPython, presumably due to dynamic class creation, but relatively similar):

```
$ python3 slots-test.py
Slots    => 0.17895982996560633
Nonslots=> 0.18887511501088738
```

For more on slots in general, see the Python standard manual set. Also, watch for the `Private` decorator case study of [Chapter 39](#)—an example that naturally allows for attributes based on both `__slots__` and `__dict__` storage, by using delegation and storage-neutral accessor tools like `getattr`.

Properties: Attribute Accessors

Our next attribute-related topic is *properties*—a mechanism that provides another way for classes to define methods called automatically for access or assignment to instance attributes. This feature is similar to “getters” and “setters” in languages like Java and C#, but in Python is generally best used sparingly as a way to add accessors to attributes *after the fact* as needs evolve and warrant. Where needed, though, properties allow attribute values to be computed dynamically without requiring method calls at the point of access.

Though properties cannot support generic attribute routing goals, at least for specific attributes, they are an alternative to some traditional uses of the `__getattr__` and `__setattr__` overloading methods we first studied in [Chapter 30](#). Properties can have a similar effect to these two methods but, by contrast, incur an extra method call only for accesses to names that require dynamic computation—other nonproperty names are accessed normally with no extra calls. Although `__getattr__` is invoked only for *undefined* names, the `__setattr__` method is instead called for assignment to *every* attribute.

Properties and slots are related, too, but serve different goals. Both implement instance attributes that are not physically stored in instance namespace dictionaries—a sort of “virtual” attribute—and both are based on the notion of class-level attribute *descriptors*. In contrast, slots manage instance storage,

while properties intercept access and compute values arbitrarily. Because their underlying descriptor implementation tool is too advanced for us to cover here, properties and descriptors both get full treatment in [Chapter 38](#).

Property basics

As a brief introduction, though, a property is a type of object assigned to a class attribute name. You can generate a property by calling the `property` built-in function, passing in up to three accessor methods—handlers for get, set, and delete operations—as well as an optional docstring for the property. If any argument is passed as `None` or omitted, that operation is not supported.

The resulting property object is typically assigned to a name at the top level of a `class` statement as a class attribute (e.g., `name=property(...)`), and a special `@` decorator syntax you'll meet later is available to automate this step. When thus assigned, later accesses to the class property name itself as an object attribute (e.g., `obj.name`) are automatically routed to one of the accessor methods passed into the `property` call.

For example, we've seen how the `__getattr__` operator-overloading method allows classes to intercept *undefined* attribute references:

```
>>> class WithOperators:
...     def __getattr__(self, name):      # On undefined ...
...         if name == 'age':
...             return 40
...         else:
...             raise AttributeError(name)
...
>>> x = WithOperators()
>>> x.age                                # Runs __getattr__
40
>>> x.name                               # Runs __getattr__
AttributeError: name
```

Here is the same example, coded with *properties* instead:

```
>>> class WithProperties:
...     def getage(self):
...         return 40
...     age = property(getage)      # (get?, set?, del?,
```

```
>>> x = WithProperties()
>>> x.age                                # Runs getage
40
>>> x.name                               # Normal fetch
AttributeError: 'WithProperties' object has no attribut
```

For some coding tasks, properties can be less complex and quicker to run than the traditional techniques. For example, when we add attribute *assignment* support, properties become more attractive—there's less code to type, and no extra method calls are incurred for assignments to attributes we don't wish to manage or compute dynamically:

```
>>> class WithProperties:
    def getage(self):
        print('get age')
        return 40
    def setage(self, value):
        print('set age:', value)
        self._age = value
    age = property(getage, setage)

>>> x = WithProperties()
>>> x.age                                # Runs getage
get age
40
>>> x.age = 42                             # Runs setage
set age: 42
>>> x._age                               # Normal fetch: no g
42
>>> x.job = 'hacker'                      # Normal assign: no s
>>> x.job                                # Normal fetch: no g
'hacker'
```

The equivalent class based on operator overloading incurs extra method calls for assignments to attributes not being managed and needs to route attribute assignments through the attribute dictionary to avoid loops (or to the `object` superclass's `__setattr__` to better support “virtual” attributes such as slots and properties coded in other classes):

```
>>> class WithOperators:
    def __getattr__(self, name):           # On ur
```

```

        if name == 'age':
            print('get age')
            return 40
        else:
            raise AttributeError(name)
    def __setattr__(self, name, value):      # On all
        print('set:', name, value)
        if name == 'age':
            self.__dict__['_age'] = value    # Or ok
        else:
            self.__dict__[name] = value

>>> x = WithOperators()
>>> x.age                         # Runs __getattr__
get age
40
>>> x.age = 41                     # Runs __setattr__
set: age 41
>>> x._age                         # Defined: no __getattr__
41
>>> x.job = 'coder'                # Runs __setattr__ again
set: job coder
>>> x.job                          # Defined: no __getattr__
'coder'

```

Properties seem like a win for this simple example. However, some applications of `__getattr__` and `__setattr__` still require more dynamic or generic interfaces than properties directly provide.

For example, the set of attributes to be managed might be unknown when a class is coded and may not even exist in a tangible form (e.g., when *delegating* arbitrary attribute references to a wrapped and embedded object generically). In such contexts, a generic attribute handler like `__getattr__` with a passed-in attribute name may be preferable. Because such generic handlers can also support simpler cases, properties may be a redundant extension—albeit one that may avoid extra calls on assignments and one that some programmers may prefer when applicable.

For more details on both options, tune in to [Chapter 38](#). As you’ll see there, it’s also possible to code properties using the `@` symbol *function decorator* syntax—a topic introduced in brief later in this chapter and an equivalent and automatic alternative to manual assignment in the class scope:

```
class WithProperties:  
    @property  
        # Coding properties in  
    def age(self):  
        # On instance.age  
  
        ...  
  
    @age.setter  
    def age(self, value):  
        # On instance.age = v  
  
        ...
```



To make sense of this decorator syntax, though, we must move ahead.

__getattribute__ and Descriptors: Attribute Implementations

To complete our attribute-tools collection, the `__getattribute__` operator-overloading method intercepts *all* attribute references, not just undefined references. This makes it more potent than its `__getattr__` cousin we used in the prior section, but also trickier to use—it’s prone to loops much like `__setattr__`, but in different ways.

For more specialized attribute interception goals, in addition to properties and operator-overloading methods, Python provides attribute *descriptors*—classes with `__get__` and `__set__` methods, assigned to class attributes and inherited by instances, that intercept read and write accesses to specific attributes. As a preview, here’s one of the simplest descriptors you’re likely to encounter:

```
>>> class AgeDesc:  
    def __get__(self, instance, owner): return 40  
    def __set__(self, instance, value): instance._age = value  
  
>>> class WithDescriptors:  
    age = AgeDesc() # Assign descriptor instance variable  
  
>>> x = WithDescriptors()  
>>> x.age # Runs AgeDesc.__get__  
40  
>>> x.age = 42 # Runs AgeDesc.__set__  
>>> x._age # Normal fetch: no AgeDesc.__get__  
42
```



Descriptors have access to state-information attributes in instances of themselves as well as their client class and are, in a sense, a more general form of properties. In fact, *properties* are a simplified way to define a specific type of descriptor—one that runs functions on access. Descriptors are also used to implement the *slots* feature we met earlier, among other Python tools, and are afforded special cases in attribute *inheritance* alluded to earlier in this chapter.

Because `__getattribute__` and descriptors are too substantial to present here, we'll defer the rest of their coverage, as well as much more on properties, to [Chapter 38](#). We'll also employ them in examples in [Chapter 39](#) and study how they factor into inheritance in [Chapter 40](#). Here, the topics tour is moving on.

Static and Class Methods

Beyond the usual methods we've been using so far, classes can define two kinds of methods called without an instance: *static* methods work roughly like simple instance-less functions inside a class no matter how they're called, and *class* methods are passed a class instead of an instance. Both are similar to tools in other languages (e.g., C++ static methods). The prior chapter's bound method coverage noted these briefly, but we'll finish the story here.

To enable these special method modes, you call built-in functions named `staticmethod` and `classmethod` within the class or invoke them with the special `@name` decoration syntax you'll meet later in this chapter. The `classmethod` call is required to enable its mode; `staticmethod` is not required for instance-less methods called only through a class name but is required if such methods are called through instances.

Why the Special Methods?

As we've seen, a class's method is normally passed an instance object in its first argument to serve as the implied subject of the method call—that's the “object” in “object-oriented programming.” Though much less common, there are two formal ways to temper this model. Before we get to the syntax, let's clarify why this might matter to you.

Sometimes, programs need to process data associated with *classes* instead of instances. Consider keeping track of the number of instances created from a class or maintaining a list of all of a class’s instances that are currently in use. This type of information and its processing are associated with the class rather than its instances. That is, the information is usually stored on the class itself and processed apart from any instance.

For such tasks, simple functions coded outside a class might suffice—because they can access class attributes through the class name, they have access to class data, and never require access to an instance. However, to better associate such code with a class and to allow such processing to be customized with inheritance as usual, it would be better to code these types of functions *inside* the class itself. To make this work, we need methods in a class that are not passed, and do not expect, a `self` instance argument.

Per the prior chapter, methods accessed through the class are *plain functions* that meet some of this need but fail if accessed through an instance: the resulting *bound method* passes an instance in calls, even if the plain function doesn’t expect one. To address, Python provides *static methods*—plain functions that are nested in a class and never expect nor receive an automatic `self` argument, regardless of how they are called. They’re optional for methods only ever accessed through classes but needed for access through instances.

Although less commonly used, Python also supports *class methods*—methods of a class that are passed a class object in their first argument instead of an instance, regardless of whether they are called through an instance or a class. Such methods can access class data through their class argument—what we’ve called `self` thus far—even if called through an instance. Normal methods, sometimes called *instance methods*, still receive a subject instance when called; static and class methods do not.

Plain-Function Methods

To demo the preceding ideas, let’s suppose that we want to use class attributes to count how many instances are generated from a class. [Example 32-5](#), `hack1.py`, makes a first attempt—its class has a counter stored as a class attribute, a constructor that bumps up the counter by one each time a new instance is created, and a method that displays the counter’s value. Remember,

class attributes are stored just once on a class and shared by all instances; storing the counter this way ensures that it effectively spans all instances.

Example 32-5. `hack1.py`

```
class Hack:  
    numInstances = 0  
    def __init__(self):  
        Hack.numInstances += 1  
    def printNumInstances():  
        print('Number of instances created:', Hack.numInstances)
```

The `printNumInstances` method is designed to process class data, not instance data—it’s about *all* the instances, not any one in particular. Because of that, we want to be able to call it without having to pass an instance. Indeed, we don’t want to *make* an instance to fetch the number of instances because this would *change* the number of instances we’re trying to fetch! In other words, we want a `self`-less “static” method.

Whether this code’s `printNumInstances` works or not, though, depends on which way you call the method—through the class or through an instance. Calls to `self`-less methods made through classes work because they produce plain functions, but calls from instances produce bound methods and fail:

```
$ python3  
>>> from hack1 import Hack  
>>> a, b, c = Hack(), Hack(), Hack()      # Make three instances  
  
>>> Hack.printNumInstances()              # Okay to call  
Number of instances created: 3  
>>> a.printNumInstances()  
TypeError: Hack.printNumInstances() takes 0 positional arguments
```

Calls to instance-less methods like `printNumInstances` made through the *class* work, but calls made through an *instance* fail because an instance is automatically passed to a method that does not have an argument to receive it. If you’re able to stick with calling `self`-less methods through classes only, you already have a static method. However, to allow `self`-less methods to

be called through instances, you need to either adopt other designs or mark such methods as special. Let's look at both options in turn.

Static Method Alternatives

Short of marking a `self`-less method as special, you can sometimes achieve similar results with different coding structures. For example, if you just want to call functions that access class members without an instance, perhaps the simplest idea is to use normal functions outside the class, not class methods.

This way, an instance isn't expected in the call. The mutation in [Example 32-6](#) illustrates.

Example 32-6. `hack2.py`

```
def printNumInstances():
    print('Number of instances created:', Hack.numInstances)

class Hack:
    numInstances = 0
    def __init__(self):
        Hack.numInstances += 1
```

Because the class name is accessible to the simple function as a global variable, this works fine. Also, note that the name of the function becomes global, but only to this single module; it will not clash with names in other files:

```
>>> import hack2 as hack
>>> a = hack.Hack()
>>> b = hack.Hack()
>>> c = hack.Hack()
>>> hack.printNumInstances()          # But function n
Number of instances created: 3      # And cannot be
>>> hack.Hack.numInstances
3
```

Prior to static methods in Python, this structure was the general prescription. Because Python already provides modules as a namespace-partitioning tool, one could argue that there's not typically any need to package functions in

classes unless they implement object behavior. Simple functions within modules like the one here do much of what instance-less class methods could and are already associated with the class because they live in the same module.

This approach, though, may be subpar. For one thing, it adds to this file's scope an extra name that is used only for processing a single class. For another, the function is not directly associated with the class by structure; in fact, its `def` could be hundreds of lines away. Worse, simple functions like this cannot be customized by inheritance since they live outside a class's namespace: subclasses cannot directly replace or extend such a function by redefining it.

We might also try to make this example work by simply using a normal method and always calling it through an instance, as usual. Unfortunately, such an approach is completely unworkable if we don't have an instance available, and making an instance changes the class data, as noted earlier. A better solution would be to somehow mark a method inside a class as never requiring an instance. The next section shows how.

Using Static and Class Methods

To designate a `self`-less method that may be called through *either* the class or its instances, classes can simply call the built-in functions `staticmethod` and `classmethod`. Both mark a function object as special—requiring no instance for the former and requiring a class argument for the latter. [Example 32-7](#) shows how.

Example 32-7. allmethods.py

```
class Methods:
    def imeth(self, x):                  # Instance method: r
        print([self, x])                # Always expects a s

    def smeth(x):                      # Static method: no
        print([x])                     # Also a plain funct

    def cmeth(cls, x):                 # Class method: gets
        print([cls, x])                # Always expects a c
```

```
smeth = staticmethod(smeth)      # Make smeth a static method
cmeth = classmethod(cmeth)      # Make cmeth a class method
```

Notice how the last two assignments in this code simply *reassign* (a.k.a. *rebind*) the method names `smeth` and `cmeth`. Attributes are created and changed by any assignment in a `class` statement, so these final assignments simply overwrite the assignments made earlier by the `def`s. As you'll see in a few moments, the special `@` decorator syntax works here as an alternative to this just as it does for properties—but makes little sense unless you first understand the assignment form here that it automates.

Technically, Python supports three kinds of class-related methods with differing argument protocols:

- *Instance methods*, passed a `self` instance object (the default)
- *Static methods*, passed no extra instance object (via `staticmethod`)
- *Class methods*, passed a class object (via `classmethod`, and inherent in metaclasses)

Moreover, simple functions in a class also serve the role of static methods without requiring any extra protocol when called through a class object only. The `allmethods.py` module illustrates all three method types, so let's expand on these in turn.

Instance methods are the normal and default case that we've used in this book so far. An instance method must always be called with an instance object. When you call it through an *instance*, Python passes the instance to the first (leftmost) argument automatically; when you call it through a *class*, you must pass along the instance manually:

```
>>> from allmethods import Methods      # Normal instance
>>> obj = Methods()                  # Callable through __call__
>>> obj.imeth(1)                   # Becomes imeth()
[<allmethods.Methods object at 0x1015a79b0>, 1]
>>> Methods.imeth(obj, 2)
[<allmethods.Methods object at 0x1015a79b0>, 2]
```

Static methods, by contrast, are called without an instance argument. Unlike simple functions outside a class, their names are local to the scopes of the

classes in which they are defined, and they may be looked up by inheritance. Instance-less functions can be called through a class normally, but using the `staticmethod` built-in allows such methods to also be called through an instance. That is, the first of the following works without the `staticmethod` in the class but the second does not:

```
>>> Methods.smeth(3)          # Static method:  
[3]  
>>> obj.smeth(4)            # No instance pc  
[4]                          # Static method:  
                             # Instance not p
```

Class methods are similar, but Python automatically passes the class (not an instance) to a class method's first (leftmost) argument, whether it is called through a class or an instance:

```
>>> Methods.cmeth(5)          # Class method:  
[<class 'allmethods.Methods'>, 5]  # Becomes cmeth(  
>>> obj.cmeth(6)            # Class method:  
[<class 'allmethods.Methods'>, 6]  # Becomes cmeth()
```

In [Chapter 40](#), you'll also find that *metaclass methods*—an advanced and technically distinct method type used in the secondary class trees of types—behave similarly to the explicitly declared class methods we're exploring here.

Counting Instances with Static Methods

Now, given these built-ins, [Example 32-8](#) codes the static method equivalent of this section's instance-counting example—it marks the method as special, so it will never be passed an instance automatically.

Example 32-8. `hack_static.py`

```
class Hack:  
    numInstances = 0                      # Use stat  
    def __init__(self):  
        Hack.numInstances += 1  
    def printNumInstances():
```

```
print('Number of instances:', Hack.numInstances)
printNumInstances = staticmethod(printNumInstances)
```

Using the static method built-in, our code now allows the `self`-less method to be called through the class or any instance of it:

```
>>> from hack_static import Hack
>>> a, b, c = Hack(), Hack(), Hack()
>>> Hack.printNumInstances()                      # Call as
Number of instances: 3
>>> a.printNumInstances()                         # Instance
Number of instances: 3
```

Compared to simply moving `printNumInstances` outside the class, as prescribed earlier, this version requires an extra `staticmethod` call (or an `@` line you'll meet ahead). However, it also localizes the function name in the class scope (so it won't clash with other names in the module); moves the function code closer to where it is used (inside the `class` statement); and allows subclasses to *customize* the static method with inheritance—a more convenient and powerful approach than importing functions from the files in which superclasses are coded. The following subclass illustrates (this continues the prior session, so the count is already 3 at the start):

```
>>> class Sub(Hack):
    def printNumInstances():                      # Override
        print('Extra stuff...')                   # But call
        Hack.printNumInstances()
    printNumInstances = staticmethod(printNumInstar

>>> a, b = Sub(), Sub()
>>> a.printNumInstances()                      # Call frst
Extra stuff...
Number of instances: 5
>>> Sub.printNumInstances()                  # Call frst
Extra stuff...
Number of instances: 5
>>> Hack.printNumInstances()                 # Call orig
Number of instances: 5
```

Moreover, classes can inherit the static method without redefining it—it is run without an instance, regardless of where it is defined in a class tree:

```
>>> class Other(Hack): pass # Inherit  
  
>>> c = Other()  
>>> c.printNumInstances()  
Number of instances: 6
```

Notice how this also bumps up the *superclass*'s instance counter because its constructor is inherited and run—a behavior that begins to encroach on the next section's subject.

Counting Instances with Class Methods

Interestingly, a *class method* can do similar work here—[Example 32-9](#) has the same behavior as the static method version listed earlier, but it uses a class method that receives the instance's class in its first argument. Rather than hardcoding the class name, the class method uses the automatically passed class object generically.

Example 32-9. hack_class.py

```
class Hack:  
    numInstances = 0 # Use class variable  
    def __init__(self):  
        Hack.numInstances += 1  
    def printNumInstances(cls):  
        print('Number of instances:', cls.numInstances)  
    printNumInstances = classmethod(printNumInstances)
```

This class is used in the same way as the prior versions, but its `printNumInstances` method receives the `Hack` class, not the instance, when called from either the class or an instance:

```
>>> from hack_class import Hack  
>>> a, b = Hack(), Hack()  
>>> a.printNumInstances() # Passes class variable  
Number of instances: 2
```

```
>>> Hack.printNumInstances() # Also passes  
Number of instances: 2
```

When using class methods, though, keep in mind that they receive the most specific (i.e., *lowest*) class of the call's subject. This has some subtle implications when trying to update class data through the passed-in class. To demo, [Example 32-10](#) subclasses to customize the same way we did for static methods in the prior section, and augments `Hack.printNumInstances` to also trace its `cls` argument.

Example 32-10. `hack_class2.py`

```
class Hack:  
    numInstances = 0 # Trace class  
    def __init__(self):  
        Hack.numInstances += 1  
    def printNumInstances(cls):  
        print('Number of instances:', cls.numInstances,  
    printNumInstances = classmethod(printNumInstances)  
  
class Sub(Hack):  
    def printNumInstances(cls): # Overrides  
        print('Extra stuff...', cls) # But calls  
        Hack.printNumInstances()  
    printNumInstances = classmethod(printNumInstances)  
  
class Other(Hack): pass # Inherit
```

Running this in a REPL reveals that the lowest class is passed in whenever a class method is run—even for subclasses that have no class methods of their own:

```
>>> from hack_class2 import Hack, Sub, Other  
>>> x = Sub()  
>>> y = Hack()  
  
>>> x.printNumInstances() # Calls Hack  
Extra stuff... <class 'hack_class2.Sub'>  
Number of instances: 2 <class 'hack_class2.Hack'>  
  
>>> Sub.printNumInstances() # Calls Sub  
Extra stuff... <class 'hack_class2.Sub'>
```

```
Number of instances: 2 <class 'hack_class2.Hack'>
```

```
>>> y.printNumInstances() # ()  
Number of instances: 2 <class 'hack_class2.Hack'>
```

In the first call here, a class method call is made through an instance of the `Sub` subclass, and Python passes the lowest class, `Sub`, to the class method. All is well in this case—since `Sub`'s redefinition of the method calls the `Hack` superclass's version explicitly, the superclass method in `Hack` receives its own class in its first argument. But watch what happens for an object that inherits the class method verbatim:

```
>>> z = Other() # ()  
>>> z.printNumInstances()  
Number of instances: 3 <class 'hack_class2.Other'>
```

This last call here passes `Other` to `Hack`'s class method. This works in this example because *fetching* the counter finds it in `Hack` by class inheritance. If this method tried to *assign* to the passed class's data, though, it would update `Other`, not `Hack`! In this specific case, `Hack` is probably better off hardcoding its own class name to update its data if it means to count instances of all its subclasses, too, rather than relying on the passed-in class argument.

Counting instances per class with class methods

In fact, because class methods always receive the *lowest* class in an instance's tree:

- *Static* methods and explicit class names may be a better solution for processing data local to a class.
- *Class* methods may be better suited to processing data that may differ for each class in a hierarchy.

Code that needs to manage *per-class* instance counters, for example, might be best off leveraging class methods. To illustrate, the top-level superclass in [Example 32-11](#) uses a class method to manage state information that varies for and is stored on each class in the tree—similar in spirit to the way instance methods manage state information that varies per class instance.

Example 32-11. hack_class3.py

```
class Hack:
    numInstances = 0
    def count(cls):                      # Per-class inst
        cls.numInstances += 1              # cls is lowest
    def __init__(self):
        self.count()                     # Passes self._
    count = classmethod(count)

class Sub(Hack):
    numInstances = 0
    def __init__(self):                  # Redefines __ir
        Hack.__init__(self)

class Other(Hack):                   # Inherits __ini
    numInstances = 0
```

When run, the `Hack` class keeps track of each of its subclasses' instances, using a counter on each subclass:

```
>>> from hack_class3 import Hack, Sub, Other
>>> x = Hack()
>>> y1, y2 = Sub(), Sub()
>>> z1, z2, z3 = Other(), Other(), Other()
>>> x.numInstances, y1.numInstances, z1.numInstances
(1, 2, 3)
>>> Hack.numInstances, Sub.numInstances, Other.numInsta
(1, 2, 3)
```

Static and class methods have additional advanced roles, which we will skip here; see other resources for more use cases. In later Python versions, though, the static and class method designations became even simpler with the advent of *function decoration* syntax—a way to apply one function to another that has roles well beyond the static method use case that was one of its initial motivations. This syntax also allows us to augment *classes*—to initialize data like the `numInstances` counter in the last example, for instance. The next section explains how.

NOTE

The methods finale: For a postscript on Python’s method types, be sure to watch for coverage of metaclass methods in [Chapter 40](#)—because these are designed to process a *class* that is an instance of a metaclass, they turn out to be very similar to the class methods defined here but require no `classmethod` declaration, and apply only to the shadowy metaclass realm previewed next.

Decorators and Metaclasses

Because the `staticmethod` and `classmethod` call technique described in the prior section initially seemed obscure to some observers, a device was eventually added to make the operation simpler. Python *decorators*—similar to the notion and syntax of annotations in Java—both address this specific need and provide a general tool for adding logic that manages functions and classes or later calls to them.

This is called a “decoration,” but in more concrete terms is really just a way to run extra processing steps at function and class definition time with explicit syntax. It comes in two flavors:

- *Function decorators:* The initial entry, augment function definitions at `def` statements. They specify operation modes for both simple functions and classes’ methods by wrapping them in an extra layer of logic implemented as another function. That function is often called a *metafunction*, though this is just terminology.
- *Class decorators:* A later extension, augment class definitions at `class` statements. They wrap classes in a similar way, adding support for management of whole objects and their interfaces instead of a single function.

We met decorators very briefly in [Chapter 19](#) in relation to simple functions, but they are more general than earlier implied: they can also be used for class methods and classes and can add nearly arbitrary logic to functions and classes that go well beyond that the static- and class-method roles used as a segue here.

For instance, *function decorators* may be used to augment functions with code that logs calls made to them, checks argument types during debugging, times

calls, and so on, and can be used to manage either functions themselves or later calls to them. In the latter mode, function decorators are similar to the *delegation* design pattern we explored in [Chapter 31](#), but they are designed to augment a specific function or method call, not an entire object interface.

Python provides a few built-in function decorators for operations, such as marking static and class methods and defining properties (as sketched earlier, the `property` built-in works as a decorator automatically), but programmers can also code arbitrary decorators of their own. Although they are not strictly tied to classes, user-defined function decorators are often coded as classes to save the original functions for later dispatch, along with other data as state information.

This proved such a useful hook that it was eventually extended—*class decorators* bring augmentation to classes, too, and are more directly tied to the class model. Like their function cohorts, class decorators may manage classes themselves or later instance-creation calls and often employ *delegation* of entire interfaces in the latter mode. As you’ll find, their roles also often overlap with *metaclasses* but are a more lightweight way to achieve some goals.

Function Decorator Basics

Syntactically, a function decorator is a sort of runtime declaration about the function that follows it. A function decorator is coded on a line by itself just before the `def` statement that defines a function or method. It consists of the `@` symbol, followed by a *metafunction*—a plain function (or other callable object) of one argument, which is passed and manages another function. The code following the `@` is usually the name of a metafunction with an optional arguments list, but as of Python 3.9, it can be any expression returning a one-argument function. Listing multiple decorators on consecutive lines allows them to nest, as you’ll see later in this book.

For example, the prior section’s methods may be coded with decorator syntax like this:

```
class C:  
    @staticmethod # Function decorati
```

```
def meth():
```

```
...
```

Internally, this syntax has the same effect as the following—passing the function through the decorator and assigning the result back to the original name:

```
class C:  
    def meth():  
        ...  
        meth = staticmethod(meth)          # Name rebinding etc
```

Decoration *rebinds* the method name to the decorator’s result. The net effect is that calling the method function’s name later actually triggers the result of its `staticmethod` decorator first. Because a decorator can return any sort of object, this allows the decorator to insert a layer of logic to be run on every later call. The decorator function is free to return either the original function itself or a new *proxy* object that saves the original function passed to the decorator to be invoked indirectly after the extra logic layer runs.

With this addition, [Example 32-12](#) is a better way to code our static method code of [Example 32-8](#).

Example 32-12. `hack_static_deco.py`

```
class Hack:  
    numInstances = 0  
    def __init__(self):  
        Hack.numInstances += 1  
  
    @staticmethod  
    def printNumInstances():  
        print('Number of instances:', Hack.numInstances)
```

Here is this example in action as before:

```
>>> from hack_static_deco import Hack  
>>> a, b, c = Hack(), Hack(), Hack()  
>>> Hack.printNumInstances()          # Calls from  
Number of instances: 3
```

```
>>> a.printNumInstances()
```

```
Number of instances: 3
```

Because they also accept and return functions, the `classmethod` and `property` built-in functions may be used as decorators in the same way—as in [Example 32-13](#), which demos all three built-in decorators previewed earlier.

Example 32-13. alldecorators.py

```
class Methods:  
    def imeth(self, x):                  # Normal instance me  
        print([self, x])  
  
    @staticmethod  
    def smeth(x):                      # Static: no instance  
        print([x])  
  
    @classmethod  
    def cmeth(cls, x):                 # Class: gets class,  
        print([cls, x])  
  
    @property                         # Property: computed  
    def name(self):  
        return 'Pat ' + self.__class__.__name__
```

Running this live in a REPL proves the point:

```
>>> from alldecorators import Methods  
>>> obj = Methods()  
>>> obj.imeth(1)  
[<alldecorators.Methods object at 0x10d2839b0>, 1]  
>>> obj.smeth(2)  
[2]  
>>> obj.cmeth(3)  
[<class 'alldecorators.Methods'>, 3]  
>>> obj.name  
'Pat Methods'
```

Bear in mind that `staticmethod` and its kin here are still built-in functions; they may be used in decoration syntax just because they take a function as an argument and return a callable to which the original function name can be

rebound. In fact, any such function can be used in this way—even user-defined functions we code ourselves, as the next section explains.

A First Look at User-Defined Function Decorators

Although Python provides a handful of built-in functions that can be used as decorators, we can also write custom decorators of our own. Because of their wide utility, we’re going to devote an entire chapter to coding decorators in the final part of this book. As a quick example, though, let’s look at a simple user-defined decorator at work.

Recall from [Chapter 30](#) that the `__call__` operator-overloading method implements a function-call interface for class instances. [Example 32-14](#) uses this to code a call proxy class that saves the decorated function in the instance and catches calls to the original name. Because this is a class, it also has state information—a counter of calls made.

Example 32-14. tracer1.py

```
class tracer:
    def __init__(self, func):                  # Remember original
        self.calls = 0
        self.func = func
    def __call__(self, *args):                 # On later calls
        self.calls += 1
        print(f'call {self.calls} to {self.func.__name_}
        return self.func(*args)

@tracer
def hack(a, b, c):                         # Same as hack =
                                            # Wrap hack in c
    return a + b + c

if __name__ == '__main__':
    print(hack(1, 2, 3))                   # Really calls t
    print(hack('a', 'b', 'c'))             # Invokes __call__
```

Because the `hack` function is run through the `tracer` decorator, when the original `hack` name is called, it actually triggers the `__call__` method in the class. This method counts and logs the call and then dispatches it to the original wrapped function. Note how the `*name` argument syntax is used to

pack and unpack the passed-in arguments; because of this, this decorator can be used to wrap any function with any number of positional arguments.

The net effect, again, is to add a layer of logic to the original `hack` function. When run, the first output line comes from the `tracer` class, and the second gives the return value of the `hack` function itself:

```
$ python3 tracer1.py
call 1 to hack
6
call 2 to hack
abc
```

Trace through this example's code for more insight. As it is, this decorator works for any function that takes positional arguments, but it does not handle *keyword* arguments and cannot decorate class-level *method* functions (in short, for methods, its `__call__` would be passed a `tracer` instance only). As you'll learn in [Part VIII](#), there are a variety of ways to code function decorators, including nested `def` statements, and some of the alternatives are better suited to methods than the version shown here.

For example, by using *nested functions* with enclosing scopes for state instead of callable class instances with attributes, function decorators often become more broadly applicable to class-level *methods* too. We'll postpone the full details on this, but [Example 32-15](#) provides a brief look at this *closure*-based coding model; it uses function attributes for counter state for portability but could also leverage variables and `nonlocal` instead.

Example 32-15. tracer2.py

```
def tracer(func):                      # Remember original
    def oncall(*args):                  # On later calls
        oncall.calls += 1
        print(f'call {oncall.calls} to {func.__name__}')
        return func(*args)
    oncall.calls = 0
    return oncall

class C:
    @tracer
    def hack(self, a, b, c): return a + b + c
```

```
if __name__ == '__main__':
    x = C()
    print(x.hack(1, 2, 3))
    print(x.hack('a', 'b', 'c'))
```

The example's output is the same as its predecessor but reflects a decorated class method; more on this later.

A First Look at Class Decorators and Metaclasses

Python later generalized decorators, allowing them to be applied to classes as well as functions. In short, *class decorators* are similar to function decorators, but they are run at the end of a `class` statement to rebind a class name to a callable. As such, they can be used to either manage classes just after they are created or insert a layer of wrapper logic to manage instances when they are later created. Symbolically, the code structure:

```
def decorator(aClass): ...

@decorator                         # Class decoration syntactic sugar
class C: ...
```

is mapped to the following equivalent:

```
def decorator(aClass): ...

class C: ...                         # Name rebinding equivalent
C = decorator(C)
```

The class decorator is free to augment the class itself or return a *proxy* object that intercepts later instance construction calls. For example, in the code of “[Counting instances per class with class methods](#)”, we could use this hook to automatically augment the classes with instance counters and any other data required:

```
def count(aClass):
    aClass.numInstances = 0
    return aClass                      # Return class itself
```

```
@count
class Hack: ...                                # Same as Hack = count

@count
class Sub(Hack): ...                          # numInstances = 0 not
```

In fact, as coded, this decorator can be applied to classes *or* functions—it happily returns the object being defined in either context after initializing the object’s attribute:

```
< ━━━━ >

@count
def hack(): pass                                # Like hack = count(hack)

@count
class Hack: pass                                # Like Hack = count(Hack)

hack.numInstances                               # Both are set to zero
Hack.numInstances
```

Though this decorator manages a function or class itself, as detailed later in this book, class decorators can also manage an object’s entire *interface* by intercepting construction calls and wrapping the new instance object in a *proxy* that deploys attribute accessor tools to intercept later requests—a multilevel coding technique we’ll use to implement class attribute privacy in [Chapter 39](#). Here’s a preview of the model:

```
def decorator(cls):                            # On @
    class Proxy:
        def __init__(self, *args):            # On init
            self.wrapped = cls(*args)
        def __getattr__(self, name):         # On attr
            return getattr(self.wrapped, name)
    return Proxy

@decorator
class C: ...                                  # Like C = decorator(C)
X = C()                                       # Makes a Proxy that wraps a C, and
```

Finally, *metaclasses*, mentioned briefly earlier in this chapter, are a similarly advanced class-based tool whose roles often intersect with those of class

decorators. They provide an alternate model, which routes the creation of a class object to a subclass of the top-level `type` class (normally), at the conclusion of a `class` statement:

```
class Meta(type):  
    def __new__(meta, classname, supers, classdict):  
        ...extra logic + class creation via type call...  
  
class C(metaclass=Meta):  
    ...my creation routed to Meta... # Like C =
```

Python calls a class's metaclass to create the new class object, passing in the data defined during the `class` statement's run; if omitted, the `metaclass` simply defaults to the `type` class we explored earlier. Abstractly speaking, here's what happens at the end of `class` statements having explicit metaclasses like the preceding:

```
classname = Meta(classname, superclasses, attributedic
```

To manage the creation or initialization of a new class object, a metaclass generally redefines the `__new__` or `__init__` method of the `type` class that intercepts this call by default. The net effect, as with class decorators, is to define code to be run automatically at class creation time. Here, this binds the class name to the result of a call to a user-defined metaclass. In fact, a metaclass need not be a class at all—a possibility we'll explore later that blurs some of the distinction between this tool and decorators and even qualifies the two as functionally equivalent in some roles.

Both schemes, class decorators and metaclasses, are free to augment a class or return an arbitrary object to replace it—a hook with almost limitless class-based customization possibilities. As you'll learn later, metaclasses may also define *methods* that process their instance classes rather than normal instances of them—a technique that's similar (if not redundant) to class methods and might be emulated by methods and data in class decorator proxies, or even a class decorator that returns a metaclass instance.

Such mind-bending concepts, however, require [Chapter 40](#)'s conceptual groundwork (and quite possibly sedation).

For More Details

Naturally, there's more to the decorator and metaclass stories than shown here.

Although they are a general mechanism whose usage may be required by some packages, coding *new* user-defined decorators and metaclasses is an advanced topic of interest primarily to tool writers, not application programmers. Because of this, this book defers additional coverage until its final and optional part:

- [Chapter 38](#) shows how to code properties using function decorator syntax in more depth.
- [Chapter 39](#) focuses on decorators, and includes more comprehensive examples.
- [Chapter 40](#) covers metaclasses, and more on the class and instance management story.

Although these chapters cover advanced topics, they'll also provide us with a chance to see Python at work in substantial examples. For now, let's move on to our final class-related topic.

The super Function

To close out this chapter, we turn to `super` : a built-in function that can be used to both reference superclass attributes implicitly without naming a superclass explicitly and route method calls coherently in multiple-inheritance trees.

So far, `super` has been called out in the sidebar “[The super Alternative](#)”, as well as multiple notes along the way, but has not appeared in code. This was by design: `super` comes with substantial complexities and downsides that make it difficult to recommend to learners. In short, it's an all-or-nothing tool with several arduous coding requirements and relies on special-case and wildly implicit semantics that run counter to Python norms.

The `super` call also tends to be abused for “Java-fication” of Python code. Newcomers with backgrounds in Java often rush to use Python's `super` simply because of its similarity to a Java tool but are unaware of its much more subtle implications in Python's multiple inheritance—until adding superclasses breaks their programs.

Nevertheless, this call has grown pervasive in Python code and merits further elaboration, especially for those opting to use it naively. Hence, this section both covers `super` usage and notes its pitfalls along the way. Ultimately, though, the merit of this call, like everything else presented in this chapter and book, is ultimately yours to decide.

The super Basics

First off, let's review the explicit alternative that's more closely in step with Python's own idioms. As we've seen in this book so far, it's always possible to reference a superclass's attributes—whether method or data—by naming their desired source class explicitly:

```
>>> class C:
    def act(self):
        print('hack')

>>> class D(C):
    def act(self):
        C.act(self)          # Name superclass explicitly
        print('code')

>>> I = D()
>>> I.act()
hack
code
```

In *single-inheritance* trees like this one, the `super` alternative seems relatively straightforward at first glance: its most common form in the following automatically selects the calling class's superclass generically and implicitly when an attribute is later fetched. An explicit call like `class.method(self)`, for example, becomes an implicit `super().method()`, as in our example:

```
>>> class C:
    def act(self):
        print('hack')

>>> class D(C):
    def act(self):
        super().act()       # Reference superclass
```

```
print('code')

>>> I = D()
>>> I.act()
hack
code
```

This works as advertised and may minimize work—you don’t need to list `self` in the call, don’t need to update the call if `D`’s superclass changes in the future, and don’t need to code long superclass names or package-import paths.

The super Details

If you study the preceding code closely, though, you’ll realize that there’s something odd going on here. The `super` call somehow knows about the class, its superclass, and the `self` instance, even though none are present in its call. The backstory involves MROs, a proxy, and an algorithm that are required reading for `super` aspirants of all kinds.

A “magic” proxy

To understand how `super` works, you first need to be fluent in the *MRO algorithm* covered in [“Multiple Inheritance and the MRO”](#)—and you should review that now if you gave it a pass. The MRO is both a firm prerequisite and nested component of `super`. Given the complexity and artificial nature of the MRO, some may rule this a first strike against `super`.

Once you’ve mastered the MRO, the simplest description of `super` is this: when used in a class method, `super` returns a *proxy* object that will locate an attribute in a class *following* that of the containing class in the MRO of the `self` instance’s class. The net effect finds an attribute in a superclass or other relative of the class containing the call.

This works as expected in single-inheritance trees because the superclass naturally follows the containing class on `self`’s MRO. Really, though, this relies on deep magic. Apart from the MRO itself, `super` works by inspecting:

- The runtime *call stack* info for the calling method’s arguments

- The `__class__` variable internally added to the `__closure__` of methods that call `super`

The combo automatically locates both the `self` argument and the class containing the `super` call and then pairs the two in a special *proxy* object that routes later attribute fetches to a superclass's version of a name.

In fact, the common no-argument `super` form is equivalent to manually passing in the class containing the `super` call, along with the `self` instance. That is, within a class's method function, the following forms work the same, though the second can be used outside a method, too, and its first argument can be `__class__` inside a method:

```
super()  
super(class-containing-the-super-call, method-self-argu
```

Both forms can be used in your code, and manual arguments may be handy in some roles. Because the second may be *harder* to code and maintain than explicit class-name references, though, Python roots out the class and instance for you behind the scenes. Here's the equivalent manual version in our example:

```
class D(C):  
    def act(self):  
        super(D, self).act()      # Works the same as sup  
        print('code')            # And D is available as
```

If that all sounds complicated and strange, it's because it is. Due to its unusual semantics, the no-argument `super` call form doesn't work at all outside the context of a class's method:

```
>>> super                  # A "magic" proxy objec  
<class 'super'>  
>>> super()                # This form has no mean  
RuntimeError: super(): no arguments
```

And where it does work, its implicit pairing of class and instance is nowhere to be found in your code:

```
>>> class E(C):
    def method(self):
        proxy = super()      # self is implicit in s
        return proxy

    >>> prx = E().method()          # The normally hidden p
    >>> prx
    <super: <class 'E'>, <E object>>

    >>> prx.act()    # Find act on MRO past hidden E, bind
Hack
```

To be sure, this call's semantics resemble nothing else in Python—it's neither a bound nor nonbound method and fills in a class and `self` even though you omit both in the call. This deviates from Python's explicit `self` policy, which holds true everywhere else. As we've seen, class methods list and use `self` explicitly to make instance references apparent. Operator overloading, including constructors, implies a `self`, but this is trivial by comparison.

By hiding the instance, `super` violates this fundamental Python idiom for a single role. While that may be comfortable to those accustomed to other OOP languages, it may also qualify as a strike two to others.

Attribute-fetch algorithm

In a *single-inheritance* tree like the preceding example, `super` is straightforward because there's just one obvious follower on the MRO—the superclass of the class containing the `super` call. In fact, in this simple case, the immediate superclass can be had from an instance at `__class__.__bases__` without applying MROs at all:

```
>>> E.__mro__
(<class '__main__.E'>, <class '__main__.C'>, <class 'object'>)
>>> E().__class__.__bases__[0]
<class '__main__.C'>
```

In the more complex class trees of *multiple inheritance*, though, you must understand `super`'s full algorithm to know what it will choose in a given tree.

Here's how this works. The proxy object that `super` returns—created “magically” from runtime info as described in the prior section—uses its saved instance and class containing the call to resolve attribute references as follows:

1. Fetch the MRO of the saved `self` instance's class, available at `self.__class__.__mro__`.
2. Scan this MRO from left to right to find the saved containing class, and skip it.
3. Search the namespace dictionaries of each remaining class in the MRO from left to right until the requested attribute is found.
4. If the attribute was found and is a method, bind it with the saved `self` instance.

This procedure is run for each attribute fetch and is wholly based on MRO ordering. It must start with `self`'s MRO because the containing class's own MRO won't apply if it has been mixed with other classes; class-tree shape and hence MRO may be arbitrary for instances made from lower classes (and may even change dynamically in rare cases).

You can't ignore these underlying mechanics except in very simple class trees. Unlike in Java, the utility of *mix-in* classes in Python makes multiple inheritance from disjoint and independent superclasses a common occurrence in realistic code. And once you add multiple superclasses, you've kicked `super` up to a whole new level.

Universal deployment

Let's illustrate with code. Suppose you've written the following classes that happily deploy `super` in simple single-inheritance mode to implicitly invoke a method one level up from `C`:

```
>>> class A:  
        def act(self): print('A')  
  
>>> class B:  
        def act(self): print('B')  
  
>>> class C(B):  
        def act(self):  
            super().act()          # super applied to a
```

```
>>> C().act()                                # Make an instance ar  
B
```

If such classes later grow to use more than one superclass, though, `super`'s effects might be surprising—it does not raise an exception when the same name appears in more than one superclass of a multiple inheritance tree but will naively pick just the *leftmost* superclass having the method being run (really, the *first* per the class tree's flattened MRO). This may or may not be the class that you want, and is completely wrong if you want both:

```
>>> class C(B, A):                          # Add an A mix-in clc  
    def act(self):  
        super().act()                      # Doesn't fail on cor  
  
>>> C().act()  
B  
  
>>> class C(A, B):  
    def act(self):  
        super().act()                      # If A is Listed firs  
  
>>> C().act()  
A
```



This silently masks a source of OOP errors so common that it shows up again in this part's "Gotchas" ahead. *Explicit* calls are one way to solve this dilemma. With explicit class names, you can choose either the left class, the right class, or both, and with substantially less drama. If you might need to be explicit later, why not use this form earlier too?

```
>>> class C(A, B):                          # Explicit form  
    def act(self):  
        A.act(self)                         # You probably want t  
        B.act(self)                         # This handles both s  
                                         # So why use the sup  
  
>>> C().act()  
A  
B
```



Technically speaking, this example's explicit *class inheritance* also searches the metaclass type tree for names not defined in superclasses, per “[The Inheritance Bifurcation](#)”. This secondary-tree search differs from `super`, which always searches just a tail portion of the *instance*'s MRO (and hence just the superclass tree), but is completely moot for defined names and unlikely to matter for undefined names.

The real underlying issue with `super` here, though, is that it requires itself to be called in *every* class's method in order to propagate the call chain. Without this *universal deployment*, the call dies in the first class that doesn't call `super`—as in our `A` and `B`. While we can add `super` to these classes, too, this is the first of that handful of arduous `super` coding requirements and quickly leads to another issue, as the next section explains.

Call-chain anchors

Since both `A` and `B` of the prior section's example are somewhere on `C`'s MRO, we might be tempted to make both classes' methods run by propagating the call with added `super` calls in both.

This scheme is called *cooperative method dispatch*: each class in a tree runs `super` to hand the call off to the next class on the MRO that cares about it. This automatically routes calls through each calling class just once and avoids running a method in a diamond's common superclass more than once (it appears just once in an MRO). It assumes that the MRO's order makes sense for your method calls, but explicit class-name calls are a fallback if not.

Armed with that info, here's the mod in our example:

```
>>> class A:  
    def act(self):  
        print('A')  
        super().act()          # Add a super here to  
  
>>> class B:  
    def act(self):  
        print('B')  
        super().act()          # Add a super here to  
  
>>> class C(B, A):  
    def act(self):
```

```
super().act()          # Hope this winds up

>>> C().act()
B
A
AttributeError: 'super' object has no attribute 'act'
```

Immediately, though, we're in trouble here, and the reason requires inspecting the MRO of an instance's class:

```
>>> I = C()
>>> I.__class__.__mro__
(<class '__main__.C'>, <class '__main__.B'>, <class '__
<-->
```

Here's the subtle problem. The `super` in `C` will select `act` in `B`, the next on this MRO; the `super` in `B` will then select `act` in `A`, its follower on `self`'s MRO; but then there are no more `act` definitions to be had: the `super` in `A` looks for `act` in `object` and beyond, and of course fails. We say that there is no *call-chain anchor*—no end point for the call propagation. Hence, the last `super` call in `A` dies with an exception.

In fact, you've just met a possible strike three for this call. While this code works if you omit the `super` in `A`, this policy won't help in general: classes like `A` and `B` are probably designed to be mixed into other classes, too, and it wouldn't make sense to specialize their code just for the `C` class's use case. To propagate `super` method calls, *all* classes must define `super`, too, and there must be an anchor to catch and end the chain somewhere.

Although we can add an anchor in a pointless superclass that defines the method but does not call `super` again, this is another of those arduous coding requirements—and substantially more effort than simply running the explicit class-name calls alternative shown earlier:

```
>>> class X:                      # Code a bogus cl
    def act(self):
        print('anchor')

>>> class A: ...same...

>>> class B: ...same...
```

```

>>> class C(B, A, X):                      # Add a final anchor
    def act(self):
        super().act()

>>> C().act()
B
A
anchor

>>> [c.__name__ for c in C().__class__.__mro__]
['C', 'B', 'A', 'X', 'object']

```

This works because `X` precedes `object` on the MRO as shown, and hence stops the `act` call chain. Adding the anchor class `X` as a common superclass to both `A` and `B` in a diamond would work, too, because the resulting MRO is the same (remember, the MRO removes all but the last [rightmost] appearance of a class from the DFLR order):

```

>>> class X: ...same...

>>> class A(X): ...same...

>>> class B(X): ...same...

>>> class C(B, A): ...same...

>>> C().act()
B
A
anchor

>>> [c.__name__ for c in C().__class__.__mro__]
['C', 'B', 'A', 'X', 'object']

```

Again, though, if you have to code a special class just to appease `super`, why not just use explicit class names? Similarity to other languages isn't a very good reason, especially in contexts that other languages don't support.

Also, keep in mind that the class selected by `super` for an attribute reference may not be a superclass at all and may vary per tree that a class is mixed into. For instance, the `super` in `B` in our example dispatches to `A`—the next one on the MRO, but a *sibling*, not a superclass. This may be moot in most programs,

but if you must be sure that an immediate superclass's method is run, you again must use explicit class names instead of `super`.

Same argument lists

While `super` always demands a call-chain anchor, you may occasionally get one for free. As a special case, `object` defines a *constructor* that can be relied on to anchor some chains (recall that the `__init__` constructor method is a class attribute like any other, despite its odd name and automatic invocation):

```
>>> class A:
    def __init__(self):
        print('A')
        super().__init__()      # Propagate constr

>>> class B:
    def __init__(self):
        print('B')
        super().__init__()      # Propagate constr

>>> class C(B, A):
    def __init__(self):
        super().__init__()      # Assume object ar

>>> I = C()
B
A
```



This propagates the constructor call through `C`, `B`, `A`, and `object` per the `I` instance's MRO via cooperative method dispatch as before. This also fails, however, for constructors that take any arguments because that of `object` takes none except `self`:

```
>>> class A:
    def __init__(self, name):      # Add an argu
        print('A')
        super().__init__(name)

>>> class B:
    def __init__(self, name):
```

```
print('B')
super().__init__(name)

>>> class C(B, A):
    def __init__(self, name):
        super().__init__(name)           # But object

>>> I = C('Pat')
B
A
TypeError: object.__init__() takes exactly one argument
```

And now you've run into another one of those other arduous coding requirements: `super` generally assumes that all the methods in a call chain use the *same arguments* list because the MRO's ordering of method calls can vary with class-tree shape: an arbitrary change in inheritance may change call order arbitrarily.

This limits flexibility inherently. While you may be able to ensure same arguments for classes used only in a single program and can sometimes fudge it with starred-argument collectors for generality, neither policy will apply to code meant to be reused in multiple contexts—which is really one of the main points behind Python programming.

Noncalls and operator overloading

To close, here are two more `super` oddities. First, keep in mind that `super`, like the MRO, is not just about methods, despite their prevalence in its jargon. It can also fetch the class *data attributes* we met in [Chapter 29](#) and the *bound methods* we met in [Chapter 31](#):

```
>>> class C:
    attr1 = 'hack'                  # super also fetches dc
    def attr2(self):                # And returns bound met
        return 'code'

>>> class D(C):
    def act(self):
        return super().attr1, super().attr2

>>> I = D()
>>> I.act()
```

```
('hack', <bound method C.attr2 of <__main__.D object at
>>> I.act()[1]()
'code'
```

When you access a method like `attr2` from a `super` proxy, the proxy binds it with the saved instance to produce a bound method. Fetching a method as a *plain function*, however, may require an explicit class name, like `C.attr2`. This leads down a rabbit hole too deep to plumb here, but it's another way that `super` clashes with normal semantics.

Second, `super` also doesn't fully work in the presence of `__X__` operator-overloading methods. If you study the following code, you'll see that explicit named calls to overloading methods in the superclass work normally, but using the `super` result in an expression fails to dispatch to the superclass's method:

```
>>> class C:
    def __getitem__(self, ix):          # Indexing over
        print('C index')

>>> class D(C):
    def __getitem__(self, ix):          # Redefine to e
        print('D index')
    C.__getitem__(self, ix)            # Explicit call
    super().__getitem__(ix)           # Direct name c
    super()[ix]                      # But operators
```

```
>>> I = C()
>>> I[99]
C index
>>> I = D()
>>> I[99]
D index
C index
C index
TypeError: 'super' object is not subscriptable
```

This behavior is due to the same limitation described in the sidebar “[Delegating Built-ins—or Not](#)”—because the proxy object returned by `super` uses the `__getattribute__` method we met earlier to catch and dispatch later attribute requests, it fails to intercept the automatic `__X__`

method invocations run by built-in operations including expressions, as these begin their search in the class instead of the instance.

This may seem less severe than the other limitations we've met, but operators should generally work the same as the equivalent method call, especially for a built-in like this. Not supporting this adds another exception for `super` users to confront and remember. Other languages' mileage may vary, but in Python, `self` is explicit, multiple-inheritance mix-ins and operator overloading are common, and superclass name changes are rare enough to pass as a red herring.

The super Wrap-Up

So there you have it: a brief tutorial on the `super` built-in, for which you can find copious supplements in all the standard places, some of which seem as focused on defending `super` as on documenting it. Hopefully, the coverage here has given you a balanced view of this tool's trade-offs while introducing its fundamentals.

As we've just seen, in single-inheritance class trees, the `super` call may be used to refer to parent superclasses generically without naming them explicitly. In multiple-inheritance trees, this call can also be used to implement cooperative method dispatch that propagates calls through a tree. The latter role may be especially useful in diamonds, as a conforming method call chain visits each superclass just once.

While these are clear upsides in some contexts, it's important to know that `super` can also yield highly implicit behavior, which for some programs may not invoke superclasses as expected or required.

To summarize, the `super` method-dispatch technique generally imposes three main coding requirements:

- *Anchors*: the method called by `super` must exist—which requires extra code and calls if no call-chain anchor is present, and mix-in classes can't be specialized for a single tree's context.
- *Arguments*: the method called by `super` must have the same argument signature across the entire class tree—which can impair flexibility, especially for implementation-level methods like constructors.

- *Deployment*: every appearance of the method called by `super` but the last must use `super` itself—which can make it difficult to use existing code, change call ordering, override methods, and code self-contained classes.

In addition, `super` builds upon the already complex MRO, can mask problems when single-inheritance trees become multiple-inheritance trees, may select a class other than a superclass in multiple-inheritance trees, and constitutes yet another special case for attribute inheritance, which we'll revisit in [Chapter 40](#).

In the end, `super` is easy to use and relatively harmless in single-inheritance roles, but its unusual semantics, rigid requirements, and questionable net reward make it a mixed bag. Python programmers, especially those learning Python anew, might be better served by the more general and transparent coding paradigm of explicit class-name references.

But you should judge all this for yourself in an OOP Python program near you.

Class Gotchas

We've reached the end of the primary OOP coverage in this book. After exceptions up next, we'll explore additional class-related examples and topics in the last part of the book, but that part mostly just gives expanded coverage to concepts introduced here. As usual, let's wrap up this part with the standard warnings about pitfalls to avoid.

Most class issues can be boiled down to namespace issues—which makes sense, given that classes are largely just namespaces with a handful of extra tricks. Some of the items in this section are more like class usage pointers than problems, but even experienced class coders have been known to stumble on a few.

Changing Class Attributes Can Have Side Effects

Theoretically speaking, classes (and class instances) are *mutable* objects. As with built-in lists and dictionaries, you can change them in place by assigning

to their attributes—and as with lists and dictionaries, this means that changing a class or instance object may impact multiple references to it.

That's usually what we want and is how objects change their state in general, but awareness of this issue becomes especially critical when changing *class* attributes. Because all instances generated from a class share the class's namespace, any changes at the class level are reflected in all instances unless they have their own versions of the changed class attributes.

In Python, we can normally change any attribute in any object to which we have a reference. Consider the following class. Inside the class body, the assignment to the name `a` generates an attribute `X.a`, which lives in the class object at runtime and will be inherited by all of `X`'s instances:

```
>>> class X:  
...     a = 1          # Class attribute  
  
>>> I = X()  
>>> I.a           # Inherited by instance  
1  
>>> X.a           # Accessible through class  
1
```

So far, so good—this is the normal case. But notice what happens when we change the class attribute dynamically *outside* the `class` statement: it also changes the attribute in every object that inherits from the class. Moreover, new instances created from the class during this session or program run also get the dynamically set value, regardless of what the class's source code says:

```
>>> X.a = 2        # May change more than X  
>>> I.a           # I changes too  
2  
>>> J = X()        # J inherits from X's runtime value  
>>> J.a           # (but assigning to J.a changes a i  
2
```



Is this a useful feature or a dangerous trap? You be the judge. As discussed in [Chapter 27](#), you can actually get work done by changing class attributes without ever making a single instance—a technique that can simulate the use

of “records” or “structs” in other languages. As a refresher, consider the following unusual but legal Python program:

```
class X: pass          # Make a few attributes
class Y: pass

X.a = 1                # Use class attributes
X.b = 2                # No instances anywhere
X.c = 3
Y.a = X.a + X.b + X.c

for X.i in range(Y.a): print(X.i)    # Prints 0..5
```



Here, the classes `X` and `Y` work like “fileless” modules—namespaces for storing variables we don’t want to clash. This is a perfectly legal Python programming trick, but it’s less appropriate when applied to classes written by others; you can’t always be sure that class attributes you change aren’t critical to the class’s internal behavior. If you’re out to simulate a C `struct`, you may be better off changing instances than classes, as that way, only one object is affected:

```
class Record: pass
X = Record()
X.name = 'pat'
X.job  = 'Pizza maker'
```

Changing Mutable Class Attributes Can Have Side Effects, Too

This gotcha is really an extension of the prior. Because class attributes are shared by all instances, if a class attribute references a *mutable* object, changing that object in place from any instance impacts all instances at once:

```
>>> class C:
        shared = []          # Class attribute
        def __init__(self):
            self.perobj = []    # Instance attribute

>>> x, y = C(), C()      # Two instances
>>> y.shared, y.perobj    # Implicitly share
```

```
([], [])
```

```
>>> x.shared.append('hack')          # Impacts y's view
>>> x.perobj.append('code')         # Impacts x's data
>>> x.shared, x.perobj
(['hack'], ['code'])

>>> y.shared, y.perobj           # y sees change made
(['hack'], [])
>>> C.shared                      # Stored on class C
['hack']
```

This effect is no different than many we've seen in this book already: mutable objects are shared by simple variables, globals are shared by functions, module-level objects are shared by multiple importers, and mutable function arguments are shared by the caller and the callee. All of these are cases of general behavior—multiple references to a mutable object—and all are impacted if the shared object is changed in place from any reference.

Here, this occurs in class attributes shared by all instances via inheritance, but it's the same phenomenon at work. It may be made more subtle by the different behavior of assignments to instance attributes themselves:

```
x.shared.append('hack')      # Changes shared object attribute
x.shared = 'hack'             # Changed or creates instance attribute
```

But again, this is not a problem, it's just something to be aware of; shared mutable class attributes can have many valid uses in Python programs.

Multiple Inheritance: Order Matters

This may be obvious by now, but it's worth underscoring one last time: if you use multiple inheritance, the order in which superclasses are listed in the `class` statement header can be critical. Python always searches superclasses from left to right, according to their order in the header line.

For instance, in the multiple inheritance example we studied in [Chapter 31](#), imagine that the `Super` class implemented a `__str__` method, too:

```
class ListTree:  
    def __str__(self): ...  
  
class Super:  
    def __str__(self): ...  
  
class Sub(ListTree, Super):      # Get ListTree's __str__  
  
x = Sub()                      # Inheritance searches L
```

Which class would we inherit it from—`ListTree` or `Super`? As inheritance searches generally proceed from left to right, we would get the method from whichever class is listed first (leftmost) in `Sub`'s `class` header. Presumably, we would list `ListTree` first because its whole purpose is its custom `__str__`. Indeed, we had to do this in [Chapter 31](#) when mixing this class with a `tkinter.Button` that had a `__str__` of its own.

But now, suppose `Super` and `ListTree` have their own versions of other same-named attributes, too. If we want one name from `Super` and another from `ListTree`, the order in which we list them in the `class` header won't help—we will have to override inheritance by manually assigning to the attribute name in the `Sub` class:

```
class ListTree:  
    def __str__(self): ...  
    def other(self): ...  
  
class Super:  
    def __str__(self): ...  
    def other(self): ...  
  
class Sub(ListTree, Super):      # Get ListTree's __str__  
    other = Super.other          # But explicitly pick Su  
    def __init__(self):  
        ...  
  
x = Sub()                      # Inheritance searches S
```

Here, the assignment to `other` within the `Sub` class creates `Sub.other`—a reference back to the `Super.other` object. Because it is lower in the

tree, `Sub.other` effectively hides `ListTree.other`, the attribute that the inheritance search would normally find. Similarly, if we listed `Super` first in the `class` header to pick up its `other`, we would need to select `ListTree`'s method explicitly:

```
class Sub(Super, ListTree):          # Get Super's  
    __str__ = ListTree.__str__        # Explicitly
```



For another example of the technique shown here in action, see the discussion of explicit conflict resolution in “[Attribute Conflict Resolution](#)”. Ultimately, multiple inheritance is an advanced tool. Even if you understood the last paragraph, it’s still a good idea to use it sparingly and carefully. Otherwise, the meaning of a name may come to depend on the order in which classes are mixed in an arbitrarily far-removed subclass.

As a rule of thumb, multiple inheritance works best when your mix-in classes are as self-contained as possible—because they may be used in a variety of contexts, they should not make assumptions about names related to other classes in a tree. The pseudoprivate `__X` attributes feature we studied in [Chapter 31](#) can help by localizing names that a class relies on owning and limiting the names that mix-in classes add to the mix. In this example, for instance, if `ListTree` only means to export its custom `__str__`, it can name its `other` method `__other` to avoid clashing with like-named classes in the tree.

Scopes in Methods and Classes

When working out the meaning of names in class-based code, it helps to remember that classes introduce local scopes, just as functions do, and methods are simply further nested functions. In the following example, the `generate` function returns an instance of the nested `Hack` class. Within its code, the class name `Hack` is assigned in the `generate` function’s local scope and hence is visible to any further nested functions, including code inside `method`; it’s in the *E* enclosing-function layer of the LEGB scope lookup rule:

```
def generate():  
    class Hack:          # Hack is a name in g  
        count = 1
```

```
def method(self):
    print(Hack.count)      # Visible in generate'
return Hack()

generate().method()
```

This example works because the local scopes of all enclosing function `def`s are automatically visible to nested `def`s—including nested method `def`s, as in this example.

Even so, keep in mind that method `def`s cannot see the local scope of the enclosing `class`; they can see only the local scopes of enclosing `def`s. That's why methods must go through the `self` instance or the class name to reference methods and other attributes defined in the enclosing `class` statement. For example, code in the method must use `self.count` or `Hack.count`, not just `count`.

To avoid nesting, we could restructure this code such that the class `Hack` is defined at the top level of the module: the nested `method` function and the top-level `generate` will then both find `Hack` in their global scopes; it's not localized to a function's scope, but is still local to a single module:

```
def generate():
    return Hack()

class Hack:                      # Define at top Level of
    count = 1
    def method(self):
        print(Hack.count)        # Found in global scope

generate().method()
```

Code tends to be simpler in general if you avoid nesting classes and functions. On the other hand, class nesting is useful in *closure* contexts, where the enclosing function's scope retains *state* used by the class or its methods. In the following, the nested `method` has access to its own scope, the enclosing function's scope (for `label`), the enclosing module's global scope, anything saved in the `self` instance by the class, and the class itself via its nonlocal name:

```
>>> def generate(label):      # Returns a class instead
    class Hack:
        count = 1
        def method(self):
            print(f'{label}={Hack.count}')
    return Hack

>>> aclass = generate('Gotchas')
>>> I = aclass()
>>> I.method()
Gotchas=1
```

Miscellaneous Class Gotchas

Here's a handful of additional class-related warnings, mostly as review:

- **Choose per-instance or class storage wisely.** On a similar note, be careful when you decide whether an attribute should be stored on a class or its instances: the former is shared by all instances, and the latter will differ per instance. In a GUI program, for instance, if you want information to be shared by all of the window class objects your application will create (e.g., the last directory used for a Save operation or an already entered password), it might be best stored as class-level data; if stored in the instance as `self` attributes, it will vary per window or be missing entirely when looked up by inheritance.
- **You usually want to call superclass constructors.** Remember that Python runs only one `__init__` constructor method when an instance is made—the first it finds by inheritance. It does not automatically run the constructors of all superclasses higher up. Because constructors normally perform required startup work, you'll usually need to run a superclass constructor from a subclass constructor—using either an explicit call through the superclass's name or `super`, passing along whatever arguments are required—unless you mean to replace the super's constructor altogether, or the superclass doesn't have or inherit a constructor at all.
- **Stay tuned for a fix for `__getattr__` and built-ins.** Another reminder: as noted in [Chapter 28](#) and elsewhere, classes that use the `__getattr__` operator-overloading method to delegate attribute fetches to wrapped objects may fail unless operator-overloading methods are

redefined in the wrapper class. The names of operator-overloading methods implicitly fetched by built-in operations are not routed through generic attribute-interception methods. To work around this, you must redefine such methods in wrapper classes, either manually, with tools, or by definition in superclasses; you'll learn how in [Chapter 39](#).

“Overwrapping-itis”

Finally, when used well, the code reuse features of OOP make it excel at cutting development time. Sometimes, though, OOP's abstraction potential can be abused to the point of making code difficult to understand. If classes are layered too deeply, code can become obscure; you may have to search through many classes to discover what an operation does.

Imagine, for example, a framework with hundreds of classes and a dozen levels of inheritance (this is a true story, but details have been omitted to protect the innocent). Deciphering method calls in such a complex system may be a monumental task: multiple classes might have to be consulted for even the most basic of operations. In fact, the logic of such a system can be so deeply wrapped that understanding a piece of code in some cases may require days of wading through related files. This obviously isn't ideal for programmer productivity.

The most general rule of thumb of Python programming applies here, too: *don't make things complicated unless they truly must be*. Wrapping your code in multiple layers of classes to the point of incomprehensibility is always a bad idea. Abstraction is the basis of polymorphism and encapsulation, and it can be a very effective tool when used well. However, you'll simplify debugging and aid maintainability if you make your class interfaces intuitive, avoid making your code overly abstract, and keep your class hierarchies short and small unless there is a good reason to do otherwise. Remember: code you write is generally code that others must read.

Chapter Summary

This chapter presented an assortment of class-related topics, including subclassing built-in types, the relationship of types and classes, slots, properties, static methods, decorators, and `super`. Most are optional

extensions to the OOP toolbox in Python but may become more useful as you start writing larger object-oriented programs, and all are fair game if they appear in code you must understand. As noted earlier, some of these topics are continued in the final part of this book; be sure to look ahead for more info on properties, descriptors, decorators, and metaclasses.

This is the end of the class part of this book, so you'll find the usual lab exercises at the end of the chapter: be sure to work through them to get some practice coding real classes. In the next chapter, we'll begin our look at our last core language topic, *exceptions*—Python's mechanism for communicating errors and other conditions to your code. This is a relatively lightweight topic but it was saved for last because new exceptions must be coded as classes. Before we tackle that final core subject, though, take a look at this chapter's quiz and the lab exercises.

Test Your Knowledge: Quiz

1. Name two ways to extend a built-in object type.
2. What are function and class decorators used for?
3. How are normal and static methods different?
4. Are tools like `__slots__` and `super` valid to use in your code?

Test Your Knowledge: Answers

1. You can embed a built-in object in a wrapper class, or subclass the built-in type directly. The latter approach tends to be simpler, as most original behavior is automatically inherited. This works because types are classes, though the way you create an instance from a type/class determines its functionality.
2. Function decorators are generally used to manage a function or method or add to it a layer of logic that is run each time the function or method is called. They can be used to log or count calls to a function, check its argument types, and so on. They are also used to “declare” static methods (simple functions in a class that are not passed an instance, however they are called), as well as class methods and properties. Class decorators are similar but manage whole objects and their interfaces instead of a function call.

3. Normal (instance) methods receive a `self` argument (the implied instance), but static methods do not. Static methods are simple functions nested in class objects. To make a method static, it must either be run through a special built-in function or be decorated with decorator syntax. Python also allows simple functions in a class to be called through the class without this step, but calls through instances still require static-method declaration.
4. *Of course*, but you shouldn't use advanced tools automatically without carefully considering their implications. Slots, for example, can break code; `super` can mask later problems when used for single inheritance, and in multiple inheritance brings with it substantial complexity for an isolated use case; and both require universal deployment to be most useful. Evaluating new or advanced tools is a primary task of any engineer, and this is why we explored trade-offs in this chapter. This book's goal is not to tell you which tools to use but to underscore the importance of objectively analyzing them—a task often given too low a priority in the software field. In engineering, as in life in general, we shouldn't let other people make choices for us.

Test Your Knowledge: Part VI Exercises

These exercises ask you to write a few classes and experiment with some existing code. Of course, the problem with existing code is that it must be existing. To work with the `set` class in exercise 5, either copy/paste

[Example 32-1](#) from emedia, find it in this book's examples package (see the [Preface](#) for pointers), or type it up by hand (mildly tedious but a great way to make syntax more concrete). These programs are growing sophisticated, so be sure to check the solutions at the end of the book for pointers. You'll find them in [Appendix B](#), under [“Part VI, Classes and OOP”](#).

1. *Inheritance*: Write a class called `Adder` that exports a method `add(self, x, y)`, which prints a “Not Implemented” message. Then, define two subclasses of `Adder` that implement the `add` method:

ListAdder

With an `add` method that returns the concatenation of its two list arguments

DictAdder

With an `add` method that returns a new dictionary containing the items in both its two dictionary arguments (any definition of dictionary addition will do; see dictionary union in [Chapter 8](#) for tips)

Experiment by making instances of all three of your classes interactively and calling their `add` methods.

Now, extend your `Adder` superclass to save an object in the instance with a constructor (e.g., assign `self.data` a list or a dictionary), and overload the `+` operator with an `__add__` method to automatically dispatch to your `add` methods (e.g., `X + Y` triggers `X.add(X.data, Y)`). Where is the best place to put the constructors and operator-overloading methods (i.e., in which classes)? What sorts of objects can you add to your class instances?

In practice, you might find it easier to code your `add` methods to accept just one real argument (e.g., `add(self, y)`) and add that one argument to the instance's current data (e.g., `self.data + y`). Does this make more sense than passing two arguments to `add`? Would you say this makes your classes more "object-oriented"?

2. *Operator overloading*: Write a class called `MyList` that shadows ("wraps") a Python list: it should overload most list operators and operations, including `+`, indexing, iteration, slicing, and list methods such as `append` and `sort`. See the Python reference manual or other documentation for a list of all possible methods to support. Also, provide a constructor for your class that takes an existing list (or a `MyList` instance) and copies its components into an instance attribute. Experiment with your class interactively. Things to explore:

- Why is copying the initial value important here?
- Can you use an empty slice (e.g., `start[:]`) to copy the initial value if it's a `MyList` instance?
- Is there a general way to route list method calls to the wrapped list?
- Can you add a `MyList` and a regular list? How about a list and a `MyList` instance?
- What type of object should operations like `+` and slicing return? What about indexing operations?
- You may implement this sort of wrapper class by embedding a real list in a standalone class or by extending the built-in list type with a subclass. Which is easier, and why?

3. *Subclassing*: Make a subclass of `MyList` from exercise 2 called `MyListSub`, which extends `MyList` to print a message to `stdout`

before each call to the `+` overloaded operation and counts the number of such calls. `MyListSub` should inherit basic method behavior from `MyList`. Adding a sequence to a `MyListSub` should print a message, increment the counter for `+` calls, and perform the superclass's method. Also, introduce a new method that prints the operation counters to `stdout` (i.e., your console window) and experiment with your class interactively. Do your counters count calls per instance or per class (for all instances of the class)? How would you program the other option? (Hint: it depends on which object the count members are assigned to: class members are shared by instances, but `self` members are per-instance data.)

4. *Attribute methods*: Write a class called `Attrs` with methods that intercept every attribute qualification (both fetches and assignments), and print messages listing their arguments to `stdout`. Create an `Attrs` instance and experiment with qualifying it interactively. What happens when you try to use the instance in expressions? Try adding, indexing, and slicing the instance of your class. (Note: a fully generic approach based upon `__getattr__` requires [Chapter 39](#)'s workarounds for reasons noted in [Chapter 28](#) and later and summarized in the solution to this exercise.)

5. *Set objects*: Experiment with the set class of [Example 32-1](#) and described in [“Extending Types by Embedding”](#). Run commands to do the following sorts of operations:

- Create two sets of integers, and compute their intersection and union by using `&` and `|` operator expressions.
- Create a set from a string, and experiment with indexing your set.
Which methods in the class are called?
- Try iterating through the items in your string set using a `for` loop.
Which methods run this time?
- Try computing the intersection and union of your string set and a simple Python string. Does it work?
- Now, extend your set by subclassing to handle arbitrarily many operands using the `*args` argument form. (Hint: see the function versions of these algorithms in [Chapter 18](#).) Compute intersections and unions of multiple operands with your set subclass. How can you intersect three or more sets, given that `&` has only two sides?
- How would you go about emulating other list operations in the set class? (Hint: `__add__` can catch concatenation, and `__getattr__`

can pass most named list method calls like `append` to the wrapped list.)

6. *Class tree links:* In “Namespaces: The Conclusion” and in “Multiple Inheritance and the MRO”, we learned that classes have a `__bases__` attribute that returns a tuple of their superclass objects (the ones listed in parentheses in the class header). Use `__bases__` to extend any or all three of the listing mix-in classes we wrote in Chapter 31 so that they print the names of the immediate superclasses of the instance’s class. Modding Example 31-10 first may be easiest. When you’re done, the first line of the string representation should look like this (your hex addresses will almost certainly vary):

```
<Instance of Sub(Super, Lister), address 0x...:>
```

7. *Composition:* Simulate a fast-food ordering scenario by defining four classes:

Lunch

A container and controller class

Customer

The actor who buys food

Employee

The actor from whom a customer orders

Food

What the customer buys

To get you started, here are the classes and methods you’ll be defining:

```
class Lunch:  
    def __init__(self)                      # Make/embed Cu  
    def order(self, foodName)                # Start a Custo  
    def result(self)                       # Ask the Custo  
  
class Customer:  
    def __init__(self)                      # Init  
    def placeOrder(self, foodName, employee) # Plac  
    def printFood(self)                    # Prin  
  
class Employee:
```

```
def takeOrder(self, foodName)      # Return a Food
```

```
class Food:
```

```
    def __init__(self, name)           # Store food na
```

The order simulation should work as follows:

- The `Lunch` class's constructor should make and embed an instance of `Customer` and an instance of `Employee`, and it should export a method called `order`. When called, this `order` method should ask the `Customer` to place an order by calling its `placeOrder` method. The `Customer`'s `placeOrder` method should, in turn, ask the `Employee` object for a new `Food` object by calling `Employee`'s `takeOrder` method.
- `Food` objects should store a food name string (e.g., “burritos”), passed down from `Lunch.order`, to `Customer.placeOrder`, to `Employee.takeOrder`, and finally to `Food`'s constructor. The top-level `Lunch` class should also export a method called `result`, which asks the customer to print the name of the food it received from the `Employee` via the `order` (this can be used to test your simulation).

Note that `Lunch` needs to pass either the `Employee` or itself to the `Customer` to allow the `Customer` to call `Employee` methods.

Experiment with your classes interactively by importing the `Lunch` class, calling its `order` method to run an interaction, and then calling its `result` method to verify that the `Customer` got what it ordered. If you prefer, you can also simply code test cases as self-test code in the file where your classes are defined, using the module `__name__` trick of [Chapter 25](#). In this simulation, the `Customer` is the active agent; how would your classes change if `Employee` were the object that initiated customer/employee interaction instead?

8. *Zoo animal hierarchy*: Consider the class tree shown in [Figure 32-1](#).

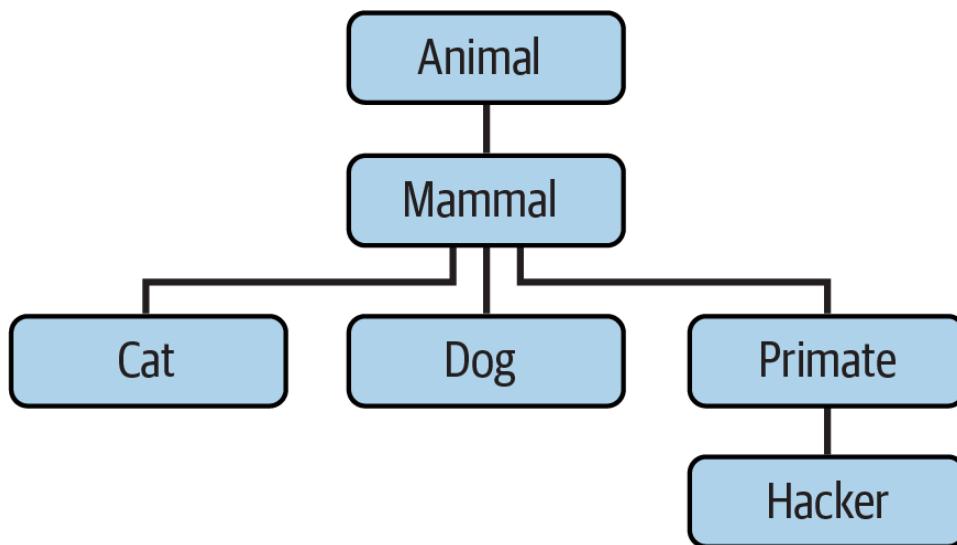


Figure 32-1. A zoo hierarchy composed of classes linked into an inheritance tree

Code a set of six `class` statements to model this taxonomy with Python *inheritance*. Then, add a `speak` method to each of your classes that prints a unique message and a `reply` method in your top-level `Animal` superclass that simply calls `self.speak` to invoke the category-specific message printer in a subclass below (this will kick off an independent inheritance search from `self`). Finally, remove the `speak` method from your `Hacker` class so that it picks up the default above it. When you’re finished, your classes should work this way:

```
$ python3
>>> from zoo import Cat, Hacker
>>> spot = Cat()
>>> spot.reply()                      # Animal.reply: c
meow
>>> data = Hacker()                  # Animal.reply: c
>>> data.reply()
Hello world!
```

WHY YOU WILL CARE: OOP BY THE MASTERS

Almost invariably, when teaching live Python classes, about halfway through the OOP section, people who have used OOP in the past are following along intensely, while people who have not are beginning to glaze over (or nod off completely). The point behind the technology just isn't apparent.

A book like this has the luxury of slowly presenting material like the overview in [Chapter 26](#), and the gradual tutorial of [Chapter 28](#)—in fact, you should probably review those sections again if you're starting to feel like OOP is just some computer science mumbo-jumbo. Though OOP adds more structure than the generators we met earlier, it similarly relies on some magic (inheritance search and a special first argument) that beginners can understandably find difficult to rationalize.

In real classes, however, to help get the newcomers on board (and keep them awake), it often helps to stop and ask the experts in the audience why they use OOP. The answers they've given might help shed some light on the purpose of OOP if you're new to the subject.

Here, then, with only a few embellishments, are the most common reasons to use OOP, as cited by students over the years:

Code reuse

This one's easy and is the main reason for using OOP. By supporting inheritance, classes make it natural to program by customization instead of starting each project from scratch.

Encapsulation

Wrapping up implementation details behind object interfaces insulates users of a class from code changes.

Structure

Classes provide new local scopes, which minimizes name clashes.

They also provide a natural place to write and look for implementation code and to manage object state.

Maintenance

Classes naturally promote code factoring, which allows us to minimize redundancy. Thanks to both the structure and code reuse support of classes, usually only one copy of the code needs to be changed.

Consistency

Classes and inheritance allow you to implement common interfaces and hence create a common look and feel in your code; this eases debugging, comprehension, and maintenance.

Polymorphism

This is more a property of OOP than a reason for using it, but by supporting code generality, polymorphism makes code more flexible and widely applicable and hence more reusable.

Other

And, of course, the number one reason students gave for using OOP: it looks good on a résumé! (OK, this one was added as a joke, but it is important to be familiar with OOP if you plan to work in the software field today.)

Finally, keep in mind the guidance given multiple times in this part of this book: you won't fully appreciate OOP until you've used it for a while. Pick a project, study larger examples, work through the exercises. Do whatever it takes to get your feet wet with OOP code. It's optional stuff and may even be overkill in some contexts, but it's generally worth the effort.
